

Dr. rer. nat. Sebastian Lapuschkin

* December 16, 1986 in Würzburg

Fraunhofer Institut für Nachrichtentechnik, Heinrich-Hertz-Institut, HHI
Einsteinufer 37, 10587 Berlin
<http://www.hhi.fraunhofer.de>
sebastian.lapuschkin@hhi.fraunhofer.de
+49 (30) 31002-371 • +49 (177) 483-2754



github.com/sebastian-lapuschkin • [linkedin.com/in/sebastian-lapuschkin](https://www.linkedin.com/in/sebastian-lapuschkin)
scholar.google.com/citations?user=wpLQuroAAAAJ

Short Bio

Sebastian Lapuschkin heads the Explainable Artificial Intelligence research group at the Fraunhofer Heinrich Hertz Institute (HHI) in Berlin, directing a team of more than 25 researchers. He received his Ph.D. with distinction from the Technische Universität Berlin in 2018 for foundational contributions to Explainable AI (XAI) and interpretable machine learning.

His research investigates the inner workings of modern frontier models: mechanistic interpretability of Transformers and large language models, the attribution of model behavior to training data and internal components, and the use of these insights to debug, control, and compress neural networks. Recent contributions include real-time attribution for LLMs and

Vision Transformers (AttnLRP), attention-head-level analysis of in-context learning, and explanation-guided pruning that removes up to 95% of model parameters without loss of performance.

His foundational methods — among them Layer-wise Relevance Propagation and Concept Relevance Propagation — are widely used across academia and industry. He is a committed advocate of open science, releasing open-source toolkits such as Zennit and Quantus to support reproducible AI research, and is a recipient of the Hugo-Geiger-Prize and the 2020 Pattern Recognition Best Paper Award.

Since 2024 he co-organizes the World Conference on eXplainable Artificial Intelligence.

Professional Experience

Technological University Dublin

DUBLIN, IRELAND

External Scholar

2025 -

at the Centre of eXplainable Artificial Intelligence. The Centre is the first of its kind in the Republic of Ireland and it aims to increase further and synergise cross-centres, college and external research collaboration.

Fraunhofer Heinrich-Hertz-Institute

BERLIN, GERMANY

Contact Person

2025 -

for the Erasmus Mundus Joint Master in Intelligent Field Robotic Systems (IFRoS) associate partnership of Fraunhofer HHI.

Ethics Committee Member

2023 -

Founding member of the first ethics committee at Fraunhofer HHI.

Responsibilities: Fulfilling the task and demand of assessing ethical aspects of research on and with humans, as well as general ethical issues.

Head of Explainable Artificial Intelligence

2021 -

Research Group Leadership and direction of XAI research & applications.

(current number of staff: 1 PostDocs, 20 PhD researchers, 30+ student research assistants & 3 Technical Staff).

Research: Work towards the next generation of local-global Explainable AI approaches and XAI-based model improvement by, e.g., increasing efficiency (see also) and debugging model training, reasoning and datasets (see also). Provision of powerful modified backprop XAI for Pytorch models, and tools for reproducible XAI evaluations to the community, e.g., via Zennit, LXT and Quantus.

Further responsibilities: Project management and (funding) acquisition. Recruitment and guidance of research personnel. Founding member of the Ethics Committee at Fraunhofer HHI.

Principal Investigator

2019 - 2020

Tenured PostDoctoral researcher in the Machine Learning Group at Fraunhofer HHI.

Research: Development of Spectral Relevance Analysis, automating the detection of “Clever Hans” moments in machine learning. Measurably increasing the explanation quality of local XAI. Provision of modified backprop XAI in Keras/Tensorflow via iNNvestigate.

Further responsibilities: Project (funding) acquisition. Recruitment and guidance of PhD students and student research assistants.

Research Associate

2014 - 2018

Founding member of the Machine Learning Group at Fraunhofer HHI.

Research: Furthering XAI research with the development and evaluation of corresponding methods, as well as applications in various expert domains, resulting in several highly cited publications, open source software tools and repositories, and the first recorded encounter of the “Clever Hans” effect in machine learning via XAI.

Other contributions: Contributions to the h.266 (VVC) video codec via learnable intra-frame prediction filters. Planning and conceptualization of an HPC cluster with modern GPU hardware implemented at Fraunhofer HHI. Development and showcasing multiple XAI demos at international events.

Technische Universität Berlin

BERLIN, GERMANY

Research Associate

2013 - 2014

Research: Formalization and development of the “Layer-wise Relevance Propagation” (LRP) method of Explainable AI for explaining individual predictions of nonlinear machine learning models.

Supervision by Prof. Dr. Klaus-Robert Müller.

Student Research- & Teaching Assistant

2011 - 2013

Research: Structure and cell type detection in large histopathology images using Bag of Words features and SVM classifiers. Development of XAI for the pipeline.

Teaching: Preparation and lecturing (of exercise sessions) in the courses “Machine Learning 1” and “Machine Learning 2 – Theory and Application” and associated academic courses. Visualization and animation of data and learning algorithms discussed throughout the course work.

Research & Teaching assistant to Prof. Dr. Klaus-Robert Müller.

Student Teaching Assistant

2009 - 2011

Course instruction for algorithmic and practical foundations of computer science (B.Sc.): Basic and advanced Java development, software engineering and OOP concepts, algorithms on image and graph data, among others.

Teaching assistant to Prof. Dr. Marc Alexa, Prof. Dr. Odej Kao and Prof. Dr. Oliver Brock.

Academic Service

Explainable and Ethical AI Workshop

Program Chair

2026 -

of the 4th workshop “Explainable and Ethical AI” (XAIE4, 2026), co-located with ICPR 2026 in Lyon, France.

XAI4Science

Organizer

2024 -

of the 1st workshop “XAI4Science: From Understanding Model Behavior to Discovering New Scientific Knowledge (2025)”, co-located with ICLR 2025 at Singapore EXPO, Singapore.
TBA: 2nd of the workshop at AAAI 2026, Singapore.

Multidisciplinary Digital Publishing Institute (MDPI)

Topic Editor

2024 - 2026

for “Opportunities and Challenges in Explainable Artificial Intelligence”.
Submission pre-screening, review management and decision handling.

World Conference on eXplainable Artificial Intelligence

Steering Committee Member

2024 -

for the 3rd XAI World Conference (2025) in Istanbul, Turkey,
for the 4th XAI World Conference (2026) in Fortaleza, Brazil ,
as well as future undisclosed instances of the conference.
Panelist, Conference and Special Track co-organization.

Programme Committee Chair

2023 - 2024

for the 2nd XAI World Conference (2024) in Valetta, Malta.
Conference and Special Track co-organization.

Reviewer & PC Memberships

Reviewer, Meta-Reviewer / Area Chair, for selected venues

NeurIPS, ICML, ICLR, CVPR, ICIP, ACL, TMLR, PLoS ONE, Nature Machine Intelligence, ...

Education

Technische Universität Berlin

BERLIN, GERMANY

PhD in Machine Learning (*summa cum laude*)

2018

Date of oral defense: December 19th, 2018.

Dean's signature on Doctorate Certificate dated January 23rd, 2019.

Research and application of methods of *Explainable AI (XAI)*: Layer-wise Relevance Propagation, Deep Taylor Decomposition and Spectral Relevance Analysis.

Thesis: "Opening the machine learning black box with Layer-wise Relevance Propagation"

Supervision headed by Prof. Dr. Klaus-Robert Müller.

Master of Science in Computer Science

2013

Focus on machine learning, computer vision and large scale data analysis.

Thesis: "On Pixel-wise Predictions from Image-wise Bag of Words Classification"

Thesis supervision headed by Prof. Dr. Alexander Binder.

Bachelor of Science in Computer Science

2010

Focus on algorithms and software development

Thesis: "Keyword-Based Image Browsing of Large Image Databases"

Thesis supervision headed by Prof. Dr. Kristian Hildebrand.

Deutschhaus-Gymnasium

WÜRZBURG, GERMANY

Abitur (pre-university secondary education)

2007

Teaching

See section "Talks & Lectures / Invited Lectures" below for a list of additional invited and individual lectures held.

WS 24/25 Machine Learning Seminar.

[**Universitat de Girona / IFRoS.** Guest Lecturer. Interactive Block Seminar "An Introduction to Explainable AI", 1 full week.]

WS 23/24 Machine Learning Seminar.

[**Universitat de Girona / IFRoS.** Guest Lecturer. Interactive Block Seminar "Explainable AI", 1 full week.]

WS 23/24 Responsible Artificial Intelligence 1.

[**Technische Universität Berlin.** Co-Teaching, Lecture Design, Interactive Coding Sessions.]

WS 21/22 Machine Learning Seminar.

[**Universitat de Girona.** Guest Lecturer. Interactive Block Seminar "Explainable AI", 1 full week.]

SS 17 Seminar Cognitive Algorithms (block seminar).

[**Technische Universität Berlin.** 1:1 Student Guidance and Co-Supervision, Grading.]

WS 13/14 Python Programming for Machine Learning (block seminar).

[**Technische Universität Berlin.** Co-Teaching, Grading, Exercise Design.]

Matlab Programming for Machine Learning and Data Analysis (block seminar).

[**Technische Universität Berlin.** Co-Teaching, Grading, Exercise Design.]

SS13 Integrated Lecture Machine Learning II.

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

Python Programming for Machine Learning (block seminar).

[**Technische Universität Berlin.** Co-Teaching, Grading, Exercise Design.]

WS 12/13 Integrated Lecture Machine Learning I.

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

Matlab Programming for Machine Learning and Data Analysis (block seminar).

[**Technische Universität Berlin.** Co-Teaching, Grading, Exercise Design.]

SS12 Machine Learning II – Theory and Application.

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

Matlab Programming for Machine Learning and Data Analysis (block seminar).

[**Technische Universität Berlin.** Co-Teaching, Grading, Exercise Design.]

WS 11/12 Machine Learning I.

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]
Matlab Programming for Machine Learning and Data Analysis (block beminar).
[**Technische Universität Berlin.** Co-Teaching, Grading, Exercise Design.]

SS 11 Methodisch-praktische Grundlagen der Informatik 2 ("Algorithms and Data Structures").

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

WS 10/11 Methodisch-praktische Grundlagen der Informatik 4 ("Advanced Algorithms").

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

SS 10 Methodisch-praktische Grundlagen der Informatik 2 ("Algorithms and Data Structures").

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

WS 09/10 Methodisch-praktische Grundlagen der Informatik 4 ("Advanced Algorithms").

[**Technische Universität Berlin.** Teaching (Exercise sessions), Grading, Exercise Design.]

Talks & Lectures

Talks

excludes internal/confidential events

1. "Explainability in the Era of LLMs: New Challenges and Pathways to Actionable Insights" (2025-08-14).
2025 Workshop on Self-Supervised Learning for Signal Decoding, Aalborg, Denmark, (invited talk)
2. "Decoding AI: How Explainability Unlocks Actionable Insights" (2025-05-23).
dida conference 2025, Berlin, Germany, (invited talk)
3. "Artificial Intelligence We Can Trust – From Explainable to Actionable and Regenerative AI" (2024-02-02).
MPNE Consensus 2024 Workshop, Berlin, Germany, (invited talk)
4. "From Concepts to Prototypes – Towards Minimal Effort Post-Hoc Interpretability" (2024-01-12).
2nd Machine Teaching for XAI Workshop (MT4XAI), Valencia, Spain, (invited talk)
5. "Explaining AI with Concept Relevance Propagation" (2023-10-06).
4th Japanese-American-German Frontiers of Science (JAGFOS) Symposium, Dresden, Germany, (flash talk & poster, invited)
6. "Model-Assisted Data Analysis via XAI" (2023-07-05).
19th Machine Learning in Healthcare Meetup Berlin, Berlin Institute of Health, (invited talk)
7. "Accessing the Hidden Space with Explainable Artificial Intelligence" (2023-06-27).
Informatik-Kolloquium, Universität Bremen, (invited talk)
8. "Explainable AI and Beyond with Concept Relevance Propagation" (2023-05-24).
Data Professional Days / Data4Business Days Köln, (keynote)
9. "Beyond Heatmaps – Explaining with Concepts" (2022-10-21).
BIFOLD Graduate School Welcome Days, (invited talk)
10. "Explain to Not Forget: Defending Against Catastrophic Forgetting with XAI " (2022-08-24).
CD-MAKE 2022, (paper presentation)
11. "Zukünftige Trends in der KI und Einsatzmöglichkeiten im Bauwesen" (2022-06-24).
BIMKIT Jahresveranstaltung 2022, (keynote)
12. "Beyond Explaining" (2021-06-03).
Melanoma Patient Network Europe Meet-up – MPNE meets AI , (invited talk)
13. "Beyond Explaining: Explainable AI for Model Improvement" (2021-05-05).
Sensor and Measurement Science International 2021, (invited talk)
14. "Efficient and Effective Neural Network Pruning with Layer-wise Relevance Propagation" (2020-11-12).
Machine Learning Seminar at Fraunhofer HHI / Technische Universität Berlin
15. "Towards Best Practice in Explaining Neural Network Decisions with LRP" (2020-07-21).
IEEE World Congress on Computational Intelligence 2020 / IJCNN 2020
16. "XAI for Analyzing and Unlearning Spurious Correlations in ImageNet" (2020-07-18).
XXAI: Extending Explainable AI Beyond Deep Models and Classifiers, (ICML 2020 Workshop)
17. "XAI via LRP and SpRAY" (2020-07-02).
Ada Day at Ada Lovelace Center / Fraunhofer IIS, (invited talk)
18. "Interpretable Machine Learning through Layer-wise Relevance Propagation" (2020-02-18).
Fraunhofer Symposium Netzwert 2020

19. "Interpretable Machine Learning through Layer-wise Relevance Propagation" (2019-12-12).
Gesellschaft von Freunden des HHI e.V.
20. "Explainable Artificial Intelligence — Opening the Machine Learning Black Box with Layer-wise Relevance Propagation" (2019-09-26).
AMA Wissenschaftsrat 2019, (invited talk)
21. "Finding Clever Hans" (2019-07-16).
Universität Bamberg, (invited talk & press interview)
22. "AI – Opening the Black Box" (2019-02-25).
Robert Koch Institut, (invited talk)
23. "AI – Opening the Black Box" (2019-02-22).
Technology Innovation Day – 91 Years HHI
24. "Understanding and Comparing Deep Neural Networks for Age and Gender Classification" (2017-10-27).
ICCV'17 Workshop on Analysis and Modeling of Faces and Gestures
25. "Layer-wise Relevance Propagation" (2014-09-10).
IDA Retreat'14

Invited Lectures

Individual Lectures as Parts of Seminars and Workshops

1. "XAI as a Tool Beyond Model Understanding – From Heatmaps to Concepts and XAI Automation" (2024-11-27).
CBS CoCoNUT | Max Planck Institute for Human Cognitive and Brain Sciences | Leipzig
2. "Human-Understandable Explanations through Concept Relevance Propagation" (2023-01-12).
Machine Teaching for Humans Workshop, Madeira | University of Bergen, (invited, keynote)
3. "Towards Human-understandable Explanations with XAI 2.0" (2022-10-24).
AI4Good webinar series of the International Telecommunication Union (ITU), (streaming link)
4. "Towards Actionable XAI" (2022-09-27).
International Artificial Intelligence Doctoral Academy, (link to slides and video)
5. "Recent Advances in Explainable AI" (2022-09-08).
BB-KI-Chips Summer School Potsdam | Universität Potsdam
6. "XAI BEYOND EXPLAINING: Using Explainability for Improving Deep Machine Learning Models" (2021-08-27).
2nd Summer School on Machine Learning in Bioinformatics | Higher School of Economics Moscow, (link to video)
7. "Neuronale Netze mit LRP (richtig) erklären" (2020-08).
KI-Campus | Die Lernplattform für Künstliche Intelligenz
8. "Explainable Artificial Intelligence — Opening the Machine Learning Black Box with Layer-wise Relevance Propagation" (2019-08-16).
SIMULA Summer School on Smart cities for a Sustainable Energy Future - From Design to Practice

Supervision & Guidance

Collaboration with and (co-)supervision of 3 (Senior) Technical Staff, 2 PostDocs, 32 PhD Students, 40+ Master's Students, 3 Bachelor's Students and 9 Guest Researchers since 2017

Third-Party Funded Research Projects

xJuRAG Explainable AI für juristische RAG-Anwendungen. leading role.	2025/10 - 2028/09
EPIBMI Explainable & Physics-Informed-based Model Improvements. leading role.	2025/08 - 2026/10
ACHILLES Human-Centred Machine Learning: Lighter, Clearer, Safer. leading role.	2024/12 - 2028/11
TEMA Trusted Extremely Precise Mapping and Prediction for Emergency Management. leading role.	2022/12 - 2026/11
DAKI-FWS Data- and AI-supported Early Warning System	2021/12 - 2024/22
iToBoS Intelligent Total Body Scanner. leading role.	2021/04 - 2025/03

BerDiBa Berlin Digital Rail Operations	2021/01 - 2024/07
Patho234 Machine Learning-driven Multidimensional Imaging Analysis of Reactive and Neoplastic Lymph Nodes	2020/01 - 2022/12
TraMeExCo Transparent Medical Expert Companion	2018/09 - 2021/08

Honors & Awards

Digital Justice Award For the research project “xJuRAG: Explainable AI für juristische RAG-Anwendungen”	2025
Machine Learning and the Physical Sciences Reproducibility Badge For the paper “PINNfluence: Influence Functions for Physics-Informed Neural Networks”	2024
Stanford Top 2% Scientist Worldwide* Among the 2% most impactful researchers of 2022 ^(ranked 195,784th) – 2025 ^(ranked 97,340th) *) which is to be taken with a grain of salt.	2022 - 2025
Best Short Paper Award @ ETRA The ACM Symposium for Eye Tracking Research and Applications	2023
Pattern Recognition Best Paper Award and Pattern Recognition Medal For the paper “Explaining NonLinear Classification Decisions with Deep Taylor Decomposition”	2020
Hugo-Geiger-Prize (1st place) Förderpreis für herausragende Promotionsleistungen	2019
Freunde des HHI Nachwuchspreis Förderpreis für exzellente wissenschaftliche Arbeiten am HHI	2019
ERCIM Cor van Baayen Award (finalist) Cor Baayen Young Researcher Award	2019
Best Paper Award @ ICMLW ICML'16 Workshop on Visualization for Deep Learning	2016

Patents

Determining a Solution of a Partial Differential Equation TBA “Determining a Solution of a Partial Differential Equation”	2025 published 2025-11-03
Output Token Selection in Machine Learning Models TBA “Output Token Selection in Machine Learning Models”	2025 published 2025-09-25
Attention Head Functionalities in Machine Learning Models TBA “Attention Head Functionalities in Machine Learning Models”	2025 published 2025-05-15
Relevance Score Assignment dealing with an Attention Module and Applications thereof WO 2025163210 A1 “Relevance Score Assignment dealing with an Attention Module and Applications thereof”	2024 published 2025-02-03
A Concept Representation of a Machine Learning Model TBA “A concept representation of a machine learning model”	2024 published 2024-12-11
Analyzing an Inference of a Machine Learning Predictor WO 2023237560 A1 “Analyzing an Inference of a Machine Learning Predictor”	2023 published 2023-12-14
Method and System for Simulating an Optical Image of a Photonic and/or Electronic Device EP 4001902 A1 “Method and System for Simulating an Optical Image of a Photonic and/or Electronic Device”	2022 published 2022-05-25
Pruning and/or Quantizing Machine Learning Predictors EP 3991102 A1 “Pruning and/or Quantizing Machine Learning Predictors” US 2022/0114455 A1 “Pruning and/or Quantizing Machine Learning Predictors” WO 2020/260656 A1 “Pruning and/or Quantizing Machine Learning Predictors”	2020 published 2022-05-04 published 2022-04-14 published 2020-12-30

Relevance Score Assignment for Artificial Neural Networks


2016

US 20180018553 "Relevance Score Assignment for Artificial Neural Networks"	granted 2024-10-04
CN 107636693 "Relevance Score Assignment for Artificial Neural Networks"	granted 2022-01-11
EP 3271863 "Relevance Score Assignment for Artificial Neural Network"	granted 2021-07-28
JP 6725547 "Relevance Score Assignment for Artificial Neural Networks"	granted 2020-07-22
KR 102130162 "Assignment of Relevance Scores for Artificial Neural Networks"	granted 2020-07-06
CA 2979579 "Relevance Score Assignment for Artificial Neural Networks"	granted 2020-02-18
RU 2703343 "Relevancy Assessment for Artificial Neural Networks"	granted 2019-10-16
BR 112017019821 "Relevance Score Assignment for Artificial Neural Networks"	published 2018-05-15
WO 2016150472 A1 "Relevance score assignment for artificial neural network"	published 2016-09-29

Publications

Summary of Scientific Impact

	All	Since 2021
# Publications	111	85
# Citations	23735	19702
h-index	40	39
i10-index	68	66



per Google Scholar, retrieved on June 5th, 2026.

List of Publications

Journal Articles

- Anders C J, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2026).
"Software for Dataset-wide XAI: From Local Explanations to Global Insights with Zennit, CoRelAy, and ViRelAy".
In: *PLoS ONE* 21(1):e0336683.
<https://github.com/chr5tphr/zennit> | <https://github.com/virelay/corelay> | <https://github.com/virelay/virelay>
- Cantú E D, Wittmann R K, Abdeen O, Wagner P, Samek W, Baier M and **Lapuschkin S** (2026).
"Deep Learning-based Multi Project InP Wafer Simulation towards Unsupervised Surface Defect Detection".
In: *Journal of the Franklin Institute* 363(2):108316
- Yolcu G Ü, Weckbecker M, Wiegand T, Samek W and **Lapuschkin S** (2025).
"Sparse, Efficient and Explainable Data Attribution with DualXDA".
In: *Transactions on Machine Learning Research* qfx81N884A.
<https://github.com/gumityolcu/DualXDA>
- Dreyer M, Berend J, Labarta T, Vielhaben J, Wiegand T, **Lapuschkin S** and Samek W (2025).
"Mechanistic Understanding and Validation of Large AI Models with SemanticLens".
In: *Nature Machine Intelligence* 1–14.
<https://github.com/jim-berend/semanticlens> | Demo: <https://semanticlens.hhi-research-insights.eu>
- Pahde F, Wiegand T, **Lapuschkin S** and Samek W (2025).
"Ensuring Medical AI Safety: Explainable AI-Driven Detection and Mitigation of Spurious Model Behavior and Associated Data".
In: *Machine Learning* 114(9):206.
<https://github.com/frederikpahde/medical-ai-safety>
- Ma J, Weicken E, Pahde F, Weitz K, **Lapuschkin S**, Samek W and Wiegand T (2025).
"Künstliche Intelligenz auf dem Prüfstand: Anforderungen, Qualitätskriterien und Prüfwerkzeuge für medizinische Anwendungen [*Artificial intelligence under scrutiny: requirements, quality criteria, and testing tools for medical applications*]".
In: *Bundesgesundheitsblatt – Gesundheitsforschung – Gesundheitsschutz* 68:915-923
- Storås A M, Dreyer M, Pahde F, **Lapuschkin S**, Samek W, Halvorsen P, de Lange T, Mori Y, Hann A, Berzin T M, Parasa S and Riegler M A (2025).
"Exploring the Clinical Value of Concept-based AI Explanations in Gastrointestinal Disease Detection".
In: *Scientific Reports* 15(1):28860.
<https://github.com/AndreaStoraas/conceptXAI-GItract>

8. Weber L, Berend J, Weckbecker M, Binder A, Wiegand T, Samek W and **Lapuschkin S** (2025).
“Efficient and Flexible Neural Network Training through Layer-wise Feedback Propagation”.
In: *Transactions on Machine Learning Research* 9oToxYVOSW.
<https://github.com/leanderweber/layerwise-feedback-propagation>
9. Hedström A, Bommer P L, Burns T F, **Lapuschkin S**, Samek W and Höhne M-C M (2025).
“Evaluating Interpretable Methods via Geometric Alignment of Functional Distortions”.
In: *Transactions on Machine Learning Research* ukLxqA8zXj.
<https://github.com/annahedstroem/GEF/> | *TMLR Survey Certification*
10. Bley F, **Lapuschkin S**, Samek W and Montavon G (2025).
“Explaining Predictive Uncertainty by Exposing Second-Order Effects”.
In: *Pattern Recognition* 160:111171.
<https://github.com/florianbley/XAI-2ndOrderUncertainty>
11. Vielhaben J, **Lapuschkin S**, Montavon G and Samek W (2024).
“Explainable AI for Time Series via Virtual Inspection Layers”.
In: *Pattern Recognition* 150:110309.
<https://github.com/jvielhaben/DFT-LRP>
12. Becker S, Vielhaben J, Ackermann M, Müller K-R, **Lapuschkin S** and Samek W (2024).
“AudioMNIST: Exploring Explainable Artificial Intelligence for Audio Analysis on a Simple Benchmark”.
In: *Journal of the Franklin Institute* 361(1):418–428.
<https://github.com/soerenab/AudioMNIST>
13. Achibat R, Dreyer M, Eisenbraun I, Bosse S, Wiegand T, Samek W and **Lapuschkin S** (2023).
“From attribution maps to human-understandable explanations through Concept Relevance Propagation”.
In: *Nature Machine Intelligence* 5(9):1006–1019.
<https://github.com/rachtibat/zennit-crp> | <https://github.com/maxdreyer/crp-human-study>
14. Hedström A, Bommer P, Wickstrøm K K, Samek W, **Lapuschkin S** and Höhne M-C M (2023).
“The Meta-Evaluation Problem in Explainable AI: Identifying Reliable Estimators with MetaQuantus”.
In: *Transactions on Machine Learning Research* j3FK00HyfU.
<https://github.com/annahedstroem/MetaQuantus>
15. Weber L, **Lapuschkin S**, Binder A and Samek W (2023).
“Beyond Explaining: Opportunities and Challenges of XAI-Based Model Improvement”.
In: *Information Fusion* 92:154–176
16. Hedström A, Weber L, Krakowczyk D G, Bareeva D, Motzkus F, Samek W, **Lapuschkin S** and Höhne M-C M (2023).
“Quantus: An Explainable AI Toolkit for Responsible Evaluation of Neural Network Explanations and Beyond”.
In: *Journal of Machine Learning Research* 24(34):1–11.
<https://github.com/understandable-machine-intelligence-lab/quantus>
17. Hofmann S M, Beyer F, **Lapuschkin S**, Goltermann O, Loeffler M, Müller K-R, Villringer A, Samek W and Witte A V (2022).
“Towards the Interpretability of Deep Learning Models for Multi-modal Neuroimaging: Finding Structural Changes of the Ageing Brain”.
In: *NeuroImage* 261:119504
18. Ma J, Schneider L, **Lapuschkin S**, Achibat R, Duchrau M, Krois J, Schwendicke F and Samek W (2022).
“Towards Trustworthy AI in Dentistry”.
In: *Journal of Dental Research* 00220345221106086
19. Rieckmann A, Dworzynski P, Arras L, **Lapuschkin S**, Samek W, Onyebuchi A A, Rod N H, Ekstrøm C T (2022).
“Causes of Outcome Learning: A Causal Inference-inspired Machine Learning Approach to Disentangling Common Combinations of Potential Causes of a Health Outcome”.
In: *International Journal of Epidemiology* dyac078.
<https://github.com/ekstroem/cool> | <https://www.causesofoutcomelearning.org>
20. Slijepcevic D, Horst F, **Lapuschkin S**, Horsak B, Raberger A-M, Kranzl A, Samek W, Breiteneder C, Schöllhorn W I and Zeppelzauer M (2022).
“Explaining Machine Learning Models for Clinical Gait Analysis”.
In: *ACM Transactions on Computing for Healthcare* 3(2):14:1–27.
<https://github.com/sebastian-lapuschkin/explaining-deep-clinical-gait-classification>
21. Anders C J, Weber L, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2022).
“Finding and Removing Clever Hans: Using Explanation Methods to Debug and Improve Deep Models”.
In: *Information Fusion* 77:261–295

22. Sun J, **Lapuschkin S**, Samek W and Binder A (2022).
“Explain and Improve: LRP-inference Fine-tuning for Image Captioning Models”.
In: *Information Fusion* 77:233–246
23. Samek W, Montavon G, **Lapuschkin S**, Anders C J, and Müller K-R (2021).
“Explaining Deep Neural Networks and Beyond: A Review of Methods and Applications”.
In: *Proceedings of the IEEE* 109(3):247–278
24. Yeom S-K, Seegerer P, **Lapuschkin S**, Binder A, Wiedemann S, Müller K-R and Samek W (2021).
“Pruning by Explaining: A Novel Criterion for Deep Neural Network Pruning”.
In: *Pattern Recognition* 115:107899.
https://github.com/seulkiyeom/LRP_pruning | https://github.com/seulkiyeom/LRP_Pruning_toy_example
25. Aeles J, Horst F, **Lapuschkin S**, Lacourpaille L, and Hug F (2021).
“Revealing the Unique Features of Each Individual’s Muscle Activation Signatures”.
In: *Journal of the Royal Society Interface* 18(174):20200770.
<https://github.com/sebastian-lapuschkin/interpretable-emg-signatures>
26. Horst F, Slijepcevic D, Zeppelzauer M, Raberger AM, **Lapuschkin S**, Samek W, Schöllhorn WI, Breiteneder C, and Horsak B (2020).
“Explaining Automated Gender Classification of Human Gait”.
In: *Gait & Posture* 81(S1):159–160
27. Hägele M, Seegerer P, **Lapuschkin S**, Bockmayr M, Samek W, Klauschen F, Müller K-R and Binder A (2020).
“Resolving Challenges in Deep Learning-based Analyses of Histopathological Images using Explanation Methods”.
In: *Scientific Reports* 10:6423
28. Alber M, **Lapuschkin S**, Seegerer P, Hägele M, Schütt K T, Montavon G, Samek W, Müller K-R, Dähne S and Kindermans P-J (2019).
“iNNvestigate Neural Networks!”.
In: *Journal of Machine Learning Research* 20(93):1–8.
<https://github.com/albermax/innvestigate>
29. **Lapuschkin S**, Wäldchen S, Binder A, Montavon G, Samek W and Müller K-R (2019).
“Unmasking Clever Hans Predictors and Assessing what Machines Really Learn”.
In: *Nature Communications* 10:1069
30. Horst F, **Lapuschkin S**, Samek W, Müller K-R and Schöllhorn W I (2019).
“Explaining the Unique Nature of Individual Gait Patterns with Deep Learning”.
In: *Scientific Reports* 9:2391.
<https://github.com/sebastian-lapuschkin/interpretable-deep-gait>
31. Montavon G, **Lapuschkin S**, Binder A, Samek W and Müller K-R (2017).
“Explaining NonLinear Classification Decisions with Deep Taylor Decomposition”.
In: *Pattern Recognition* 65:211–222.
Pattern Recognition Best Paper Award and Pattern Recognition Medal winner
32. Samek W, Binder A, Montavon G, **Lapuschkin S**, and Müller K-R (2017).
“Evaluating the Visualization of what a Deep Neural Network has Learned”.
In: *IEEE Transactions of Neural Networks and Learning Systems*
33. Sturm I, **Lapuschkin S**, Samek W and Müller K-R (2016).
“Interpretable Deep Neural Networks for Single-Trial EEG Classification”.
In: *Journal of Neuroscience Methods* 274:141–145
34. **Lapuschkin S**, Binder A, Montavon G, Müller K-R and Samek W (2016).
“The Layer-wise Relevance Propagation Toolbox for Artificial Neural Networks”.
In: *Journal of Machine Learning Research* 17(114):1–5.
https://github.com/sebastian-lapuschkin/lrp_toolbox
35. **Bach S**, Binder A, Montavon G, Klauschen F, Müller K-R and Samek W (2015).
“On Pixel-wise Explanations for Non-Linear Classifier Decisions by Layer-wise Relevance Propagation”.
In: *PLoS ONE* 10(7):e0130140

Contributions to Conference Proceedings and Workshops

1. Ostermann S, Gurgurov D, Baeumel T, Hedderich M A, **Lapuschkin S**, Samek W and Schmitt V (2026).
“From Weights to Activations: Is Steering the Next Frontier of Adaptation?”.
In: *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*
2. Komorowski P, Golimblevskaia E, Achtibat R, Wiegand T, **Lapuschkin S** and Samek W (2026).
“Attribution-Guided Decoding”.

- In: *Proceedings of the International Conference on Learning Representations (ICLR)* .
<https://github.com/piotr-komorowski/attribution-guided-decoding>
3. Kahardipraja P, Achibat R, Wiegand T, Samek W and **Lapuschkin S** (2025).
 “The Atlas of In-Context Learning: How Attention Heads Shape In-Context Retrieval Augmentation”.
 In: *Advances in Neural Information Processing Systems (NeurIPS)* .
<https://github.com/pkhdipraja/in-context-atlas>
 4. Bareeva D, Höhne M M C, Warnecke A, Pirch L, Müller K-R, Rieck K, **Lapuschkin S** and Bykov K (2025).
 “Manipulating Feature Visualizations with Gradient Slingshots”.
 In: *Advances in Neural Information Processing Systems (NeurIPS)* .
https://github.com/dilyabareeva/grad_slingshot
 5. Sandmann E, **Lapuschkin S** and Samek W (2025).
 “The Algorithm Is Not the Behavior: Learned Priors Override Look-Ahead in a Chess-Playing Neural Network”.
 In: *NeurIPS 2025 Workshop on Mechanistic Interpretability* .
<https://github.com/hartigel/leela-logit-lens>
 6. Golimblevskaia E, Jain A, Puri B, Ibrahim A, Samek W and **Lapuschkin S** (2025).
 “Circuit Insights: Towards Interpretability Beyond Activations”.
 In: *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP (EMNLP)* .
<https://github.com/egolimblevskaia/WeightLens> | <https://github.com/egolimblevskaia/CircuitLens>
 7. Labarta T, Hoang N, Weitz K, Samek W, **Lapuschkin S** and Weber L (2025).
 “See What I Mean? CUE: A Cognitive Model of Understanding Explanations”.
 In: *Proceedings of the IJCAI Workshops 2025: XAI Workshop* .
<https://arxiv.org/abs/2506.14775>
 8. Puri B, Jain A, Golimblevskaia E, Kahardipraja P, Wiegand T, Samek W and **Lapuschkin S** (2025).
 “FADE: Why Bad Descriptions Happen to Good Features”.
 In: *Findings of the Association for Computational Linguistics (ACL)* 17138–17160.
<https://github.com/brunibrun/FADE>
 9. Naujoks J, Krasowski A, Weckbecker M, Yolcu G Ü, Wiegand T, **Lapuschkin S**, Samek W and Klausen R P (2025).
 “Leveraging Influence Functions for Resampling Data in Physics-Informed Neural Networks”.
 In: *Proceedings of the 3rd XAI World Conference* 383–395.
https://github.com/aleks-krasowski/PINNfluence_resampling
 10. Erogullari E, **Lapuschkin S**, Samek W and Pahde F (2025).
 “Post-Hoc Concept Disentanglement: From Correlated to Isolated Concept Representations”.
 In: *Proceedings of the 3rd XAI World Conference* 68–89.
<https://github.com/erenerogullari/cav-disentanglement>
 11. Joseph S, Suresh P, Hufe L, Stevinson E, Graham R, Vadi Y, Bzdok D, **Lapuschkin S**, Sharkey L and Richards A (2025).
 “Prisma: An Open Source Toolkit for Mechanistic Interpretability in Vision and Video”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops: MIV Workshop* TBA.
<https://arxiv.org/abs/2504.19475> | <https://github.com/Prisma-Multimodal/ViT-Prisma>
 12. Pahde F, Dreyer M, Weckbecker M, Weber L, Anders C J, Wiegand T, Samek W and **Lapuschkin S** (2025).
 “Navigating Neural Space: Revisiting Concept Activation Vectors to Overcome Directional Divergence”.
 In: *Proceedings of the International Conference on Learning Representations (ICLR)* .
<https://github.com/frederikpahde/pattern-cav>
 13. Bareeva D, Yolcu GÜ, Hedström A, Wiegand T, Samek W and **Lapuschkin S** (2024).
 “Quanda: An Interpretability Toolkit for Training Data Attribution Evaluation and Beyond”.
 In: *NeurIPS 2024 Workshop on Attributing Model Behavior at Scale (ATTRIB 2024)* .
<https://github.com/dilyabareeva/quanda>
 14. Naujoks J R, Krasowski A, Weckbecker M, Wiegand T, **Lapuschkin S**, Samek W and Klausen R P (2024).
 “PINNfluence: Influence Functions for Physics-Informed Neural Networks”.
 In: *NeurIPS 2024 Workshop on Machine Learning and the Physical Sciences (MLAPS)* .
<https://github.com/aleks-krasowski/PINNfluence>
 Reproducibility Badge Winner
 15. Kopf L, Bommer P L, Hedström A, **Lapuschkin S**, Höhne M M-C and Bykov K (2024).
 “CoSy: Evaluating Textual Explanations of Neurons”.
 In: *Advances in Neural Information Processing Systems (NeurIPS)* 34656–34685. (*OpenReview*)
<https://github.com/lkopf/cosy>

16. Nobis G, Springenberg M, Aversa M, Detzel M, Daems R, Murray-Smith R, Nakajima S, **Lapuschkin S**, Ermon S, Birdal T, Oppel M, Knochenhauer C, Oala L and Samek W (2024).
“Generative Fractional Diffusion Models”.
In: *Advances in Neural Information Processing Systems (NeurIPS)* 25469–25509. (OpenReview)
<https://github.com/GabrielNobis/gfdm>
17. Mekala R R, Pahde F, Baur S, Chandrashekar S, Diep M, Wenzel M A, Wisotzky E L, Yolcu G Ü, **Lapuschkin S**, Ma J, Eisert P, Lindvall M, Porter A and Samek W (2024).
“Synthetic Generation of Dermatoscopic Images with GAN and Closed-Form Factorization”.
In: *ECCV 2024 Workshop on Synthetic Data for Computer Vision (SyntheticData4CV)* 15642:368–384. (Green Open Access)
18. Achtabat R, Hatefi S M V, Dreyer M, Jain A, Wiegand T, **Lapuschkin S**, Samek W (2024).
“AttnLRP: Attention-Aware Layer-wise Relevance Propagation for Transformers”.
In: *Proceedings of the 41st International Conference on Machine Learning (ICML)* 135–168.
<https://github.com/rachtabat/LRP-for-Transformers>
19. Hatefi S M V, Dreyer M, Achtabat R, Wiegand T, Samek W and **Lapuschkin S** (2024).
“Pruning By Explaining Revisited: Optimizing Attribution Methods to Prune CNNs and Transformers”.
In: *Proceedings of the European Conference on Computer Vision (ECCV) Workshops* 152–169. (Green Open Access)
<https://github.com/erfanhatefi/Pruning-by-explaining-in-PyTorch>
20. Hedström A, Weber L, **Lapuschkin S**, Höhne M M-C (2024).
“A Fresh Look at Sanity Checks for Saliency Maps”.
In: *Proceedings of the 2nd XAI World Conference* 403–420. (Green Open Access)
<https://github.com/annahedstroem/sanity-checks-revisited>
21. Tinauer C, Damulina A, Sackl M, Soellradl M, Achtabat R, Dreyer M, Pahde F, **Lapuschkin S**, Schmidt R, Ropele S, Samek W, Langkammer C (2024).
“Explainable Concept Mappings of MRI: Revealing the Mechanisms Underlying Deep Learning-based Brain Disease Classification”.
In: *Proceedings of the 2nd XAI World Conference* 202–216. (Green Open Access)
22. Dreyer M, Pürelku E, Vielhaben J, Samek W, **Lapuschkin S** (2024).
“PURE: Turning Polysemantic Neurons Into Pure Features by Identifying Relevant Circuits”.
In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 8212–8217.
<https://github.com/maxdreyer/PURE | Spotlight Paper>
23. Bareeva D, Dreyer M, Pahde F, Samek W and **Lapuschkin S** (2024).
“Reactive Model Correction: Mitigating Harm to Task-Relevant Features via Conditional Bias Suppression”.
In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3532–3541.
https://github.com/dilyabareeva/reactive_correction
24. Dreyer M, Achtabat R, Samek W and **Lapuschkin S** (2024).
“Understanding the (Extra-)Ordinary: Validating Deep Model Decisions with Prototypical Concept-based Explanations”.
In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3491–3501.
<https://github.com/maxdreyer/pcx>
25. Dreyer M, Pahde F, Anders C J, Samek W and **Lapuschkin S** (2024).
“From Hope to Safety: Unlearning Biases of Deep Models via Gradient Penalization in Latent Space”.
In: *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)* 38(19):21046–21054.
<https://github.com/frederikpahde/rrclarc>
26. Dawoud K, Samek W, Eisert P, **Lapuschkin S** and Bosse S (2023).
“Human-Centered Evaluation of XAI Methods”.
In: *Proceedings of the IEEE International Conference on Data Mining (ICDM)* 912–921. (Green Open Access)
27. Frommholz A, Seipel F, **Lapuschkin S**, Samek W and Vielhaben J (2023).
“XAI-based Comparison of Audio Event Classifiers with different Input Representations”.
In: *Proceedings of the International Conference on Content-based Multimedia Indexing (CBMI)* 126–132
28. Hedström A, Weber L, **Lapuschkin S** and Höhne M M-C (2023).
“Sanity Checks Revisited: An Exploration to Repair the Model Parameter Randomisation Test”.
In: *NeurIPS 2023 Workshop on XAI (XAI in Action: Past, Present, and Future Applications)* (vVpofYmnsG)
29. Pahde F, Dreyer M, Samek W and **Lapuschkin S** (2023).
“Reveal to Revise: An Explainable AI Life Cycle for Iterative Bias Correction of Deep Models”.

- In: *Proceedings of the International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)* 596–606. (Green Open Access)
<https://github.com/maxdreyer/reveal2revise>
30. Binder A, Weber L, **Lapuschkin S**, Montavon G, Müller K-R and Samek W (2023).
 “Shortcomings of Top-Down Randomization-Based Sanity Checks for Evaluations of Deep Neural Network Explanations”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* 16143–16152
 31. Dreyer M, Achibat R, Wiegand T, Samek W and **Lapuschkin S** (2023).
 “Revealing Hidden Context Bias in Segmentation and Object Detection through Concept-specific Explanations”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3828–3838
 32. Pahde F, Yolcu GÜ, Binder A, Samek W and **Lapuschkin S** (2023).
 “Optimizing Explanations by Network Canonization and Hyperparameter Search”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3818–3827
 33. Krakowczyk D G, Prasse P, Reich D R, **Lapuschkin S**, Scheffer T, Jäger L A (2023).
 “Bridging the Gap: Gaze Events as Interpretable Concepts to Explain Deep Neural Sequence Models”.
 In: *Proceedings of the Symposium on Eye Tracking Research and Applications (ETRA)* 1–8.
Best Short Paper Award Winner
 34. Krakowczyk D G, Reich D R, Prasse P, **Lapuschkin S**, Jäger L A and Scheffer T (2022).
 “Selection of XAI Methods Matters: Evaluation of Feature Attribution Methods for Oculomotoric Biometric Identification”.
 In: *NeurIPS 2022 Workshop on Gaze Meets ML (GOLdDAP2AtI)*
 35. Motzkus F, Weber L and **Lapuschkin S** (2022).
 “Measurably Stronger Explanation Reliability via Model Canonization”.
 In: *Proceedings of the International Conference on Image Processing (ICIP)* 516–520
 36. Ede S, Baghdadian S, Weber L, Nguyen A, Zanca D, Samek W and **Lapuschkin S** (2022).
 “Explain to Not Forget: Defending Against Catastrophic Forgetting with XAI”.
 In: *Proceedings of the International Cross-Domain Conference for Machine Learning and Knowledge Extraction (CD-MAKE)* 1–18. (Gold Open Access link)
 37. Sun J, **Lapuschkin S**, Samek W, Zhao Y, Cheung N-M and Binder A (2021).
 “Explanation-Guided Training for Cross-Domain Few-Shot Classification”.
 In: *Proceedings of the 25th International Conference on Pattern Recognition (ICPR)* 7609–7616
 38. Goh G S W, **Lapuschkin S**, Weber L, Samek W and Binder A (2021).
 “Understanding Integrated Gradients with SmoothTaylor for Deep Neural Network Attribution”.
 In: *Proceedings of the 25th International Conference on Pattern Recognition (ICPR)* 4949–4956
 39. Kohlbrenner M, Bauer A, Nakajima S, Binder A, Samek W, and **Lapuschkin S** (2020).
 “Towards Best Practice in Explaining Neural Network Decisions with LRP”.
 In: *Proceedings of the IEEE International Joint Conference on Neural Networks (IJCNN)* 1-7
 40. Sun J, **Lapuschkin S**, Samek W and Binder A (2020).
 “Understanding Image Captioning Models beyond Visualizing Attention”.
 In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*
 41. Anders C J, Neumann D, Marinč T, Samek W, Müller K-R and **Lapuschkin S** (2020).
 “XAI for Analyzing and Unlearning Spurious Correlations in ImageNet”.
 In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*
 42. Sun J, **Lapuschkin S**, Samek W, Zhao Y, Cheung N-M and Binder A (2020).
 “Explain and Improve: Cross-Domain-Few-Shot-Learning Using Explanations”.
 In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*
 43. Alber M, **Lapuschkin S**, Seegerer P, Hägele M, Schütt K T, Montavon G, Samek W, Müller K-R, Dähne S and Kindermans P-J (2018).
 “How to iNNvestigate Neural Networks’ Predictors!”.
 In: *Machine Learning Open Source Software: Sustainable Communities. NIPS Workshop*
 44. **Lapuschkin S**, Binder A, Müller K-R and Samek W (2017).
 “Understanding and Comparing Deep Neural Networks for Age and Gender Classification”.
 In: *Proceedings of the ICCV’17 Workshop on Analysis and Modeling of Faces and Gestures (AMFG)* 2017:1629-1638
 45. Srinivasan V, **Lapuschkin S**, Hellge C, Müller K-R and Samek W (2017).
 “Interpretable Action Recognition in Compressed Domain”.

In: *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* 2017:1692-1696

46. **Bach S**, Binder A, Müller K-R and Samek W (2016).
"Controlling Explanatory Heatmap Resolution and Semantics via Decomposition Depth".
In: *Proceedings of the IEEE International Conference of Image Processing (ICIP)* 2016:2271-2275
47. Binder A, Samek W, Montavon G, **Bach S**, and Müller K-R (2016).
"Analyzing and Validating Neural Network Predictions".
In: *Proceedings of the ICML'16 Workshop on Visualization for Deep Learning* .
Best Paper Award Winner
48. **Lapuschkin S**, Binder A, Montavon G, Müller K-R and Samek W (2016).
"Analyzing Classifiers: Fisher Vectors and Deep Neural Networks".
In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* 2016:2912-2920
49. Montavon G, **Bach S**, Binder A, Samek W and Müller K-R (2016).
"Deep Taylor Decomposition of Neural Networks".
In: *Proceedings of the ICML'16 Workshop on Visualization for Deep Learning* 2016:1-3
50. Samek W, Montavon G, Binder A, **Lapuschkin S** and Müller K-R (2016).
"Interpreting the Predictions of Complex ML Models by Layer-wise Relevance Propagation".
In: *Proceedings of the Interpretable ML for Complex Systems NIPS'16 Workshop*

Books

1. Longo L, **Lapuschkin S** and Seifert C, editors (2024).
"Explainable Artificial Intelligence (Second World Conference, xAI 2024, Valletta, Malta, July 17–19, 2024, Proceedings, Part I-IV)".
Springer (Cham), Part I ISBN: 978-3-031-63787-2. Part II ISBN: 978-3-031-63797-1.
Part III ISBN: 978-3-031-63800-8. Part IV ISBN: 978-3-031-63803-9

Book Chapters

1. Becking D, Dreyer M, Samek W, Müller K and **Lapuschkin S** (2022).
"ECQ^{*}: Explainability-Driven Quantization for Low-Bit and Sparse DNNs".
In: *xxAI – Beyond Explainable AI* 271-296. Springer, Cham
2. Montavon G, Binder A, **Lapuschkin S**, Samek W and Müller K-R (2019).
"Layer-wise relevance propagation: An Overview".
In: *Explainable AI: Interpreting, Explaining and Visualizing Deep Learning* 193-209. Springer, Cham
3. Binder A, **Bach S**, Montavon G, Müller K-R and Samek W (2016).
"Layer-wise Relevance Propagation for Deep Neural Network Architectures".
In: *Information Science and Applications (ICISA) 2016. Lecture Notes in Electrical Engineering* 276:913-922.
Springer, Singapore
4. Binder A, Montavon G, **Lapuschkin S**, Müller K-R and Samek W (2016).
"Layer-wise Relevance Propagation for Neural Networks with Local Renormalization Layers".
In: *Lecture Notes in Computer Science* 9887:63-71. Springer, Berlin/Heidelberg

Preprints

1. Kortukov E, Komorowski P, Klein F, Engl P, Sarti G, Oh S J, **Lapuschkin S** and Samek W (2026).
"Predicting Future Behaviors in Reasoning Models Enables Better Steering".
In: *CoRR abs/2606.11172*.
Accepted at the Mechanistic Interpretability Workshop at ICML 2026
2. Pham M A, Segeler A, Wiegand T, Samek W, **Lapuschkin S**, Kahardipraja P and Achibat R (2026).
"Fast & Faithful Function Vectors".
In: *CoRR abs/2606.05079*
3. Zimmermann J P, Berend J, Loho G, **Lapuschkin S** and Samek W (2026).
"Playing the Network Backward: A Game Theoretic Attribution Framework".
In: *CoRR abs/2605.06212*
4. Bouanani O, Berend J, Samek W, **Lapuschkin S** and Dreyer M (2026).
"Contrastive Semantic Projection: Faithful Neuron Labeling with Contrastive Examples".
In: *CoRR abs/2604.22477*.
Accepted for publication at the 4th XAI World Conference (2026)
5. Labarta T, Dreyer M, Weitz K, Samek W and **Lapuschkin S** (2026).
"From Attribution to Action: A Human-Centered Application of Activation Steering".
In: *CoRR abs/2604.11467*
6. Mesgari M M, Ma J, Samek W, **Lapuschkin S** and Weber L (2026).
"Structural Compactness as a Complementary Criterion for Explanation Quality".

- In: CoRR abs/2603.29491.
Accepted for publication at the 4th XAI World Conference (2026)
7. Heydari S, Said J, Yolcu G Ü, Kortukov E, Golimblevskaia E, Vlachos E, Mygdalis V, Pitas I, **Lapuschkin S** and Arras L (2026).
“Concept-based Explanations of Segmentation and Detection Models in Natural Disaster Management”.
In: CoRR abs/2603.23020.
Accepted as a late-breaking contribution at the 4th XAI World Conference (2026)
 8. Krasowski A, Klausen R P, Celik A, **Lapuschkin S**, Samek W and Naujoks J (2026).
“Building Trust in PINNs: Error Estimation through Finite Difference Methods”.
In: CoRR abs/2603.15526.
Accepted as a late-breaking contribution at the 4th XAI World Conference (2026)
 9. Labarta T, Hoang N, Dreyer M, Berend J, Hein O, Ma J, Samek W and **Lapuschkin S** (2026).
“X-SYS: A Reference Architecture for Interactive Explanation Systems”.
In: CoRR abs/2602.12748.
Accepted for publication at the 4th XAI World Conference (2026)
 10. Puri B, Berend J, **Lapuschkin S** and Samek W (2025).
“Atlas-Alignment: Making Interpretability Transferable Across Language Models”.
In: CoRR abs/2510.27413
 11. Klausen R P, Timofeev I, Frank J, Naujoks J, Wiegand T, **Lapuschkin S** and Samek W (2025).
“LieSolver: A PDE-constrained solver for IBVPs using Lie symmetries”.
In: CoRR abs/2510.25731
 12. Panfilov A, Kortukov E, Nikolić K, Bethge M, **Lapuschkin S**, Samek W, Prabhu A, Andriushchenko M, Geiping J (2025).
“Strategic Dishonesty Can Undermine AI Safety Evaluations of Frontier LLMs”.
In: CoRR abs/2509.18058
 13. Hufe L, Venhoff C, Pürelku E, Dreyer M, **Lapuschkin S** and Samek W (2025).
“Dyslexify: A Mechanistic Defense Against Typographic Attacks in CLIP”.
In: CoRR abs/2508.20570
 14. Hatefi S M V, Dreyer M, Achtibat R, Kahardipraja P, Wiegand T, Samek W, Binder A and **Lapuschkin S** (2025).
“Attribution-Guided Pruning for Insight and Control: Circuit Discovery and Targeted Correction in Small-scale LLMs”.
In: CoRR abs/2506.13727.
<https://github.com/erfanhatefi/SparC3>
 15. Gururaj S, Grüne L, Samek W, **Lapuschkin S** and Weber L (2025).
“Relevance-driven Input Dropout: an Explanation-guided Regularization Technique”.
In: CoRR abs/2505.21595.
https://github.com/Shreyas-Gururaj/LRP_Relevance_Dropout
 16. Dreyer M, Hufe L, Berend J, Wiegand T, **Lapuschkin S** and Samek W (2025).
“From What to How: Attributing CLIP’s Latent Components Reveals Unexpected Semantic Reliance”.
In: CoRR abs/2505.20229.
<https://github.com/maxdreyer/attributing-clip>
 17. Zverev E, Kortukov E, Panfilov A, Volkova A, Tabesh S, **Lapuschkin S**, Samek W and Lampert C H (2025).
“ASIDE: Architectural Separation of Instructions and Data in Language Models”.
In: CoRR abs/2503.10566
 18. Arras L, Puri B, Kahardipraja P, **Lapuschkin S** and Samek W (2025).
“A Close Look at Decomposition-based XAI-Methods for Transformer Language Models”.
In: CoRR abs/2502.15886
 19. Gerstenberger M, **Lapuschkin S**, Eisert P and Bosse S (2022).
“But That’s Not Why: Inference Adjustment by Interactive Prototype Deselection”.
In: CoRR abs/2203.10087
 20. Anders C J, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2021).
“Software for Dataset-wide XAI: From Local Explanations to Global Insights with Zennit, CoRelAy, and ViRelAy”.
In: CoRR abs/2106.13200. <https://github.com/chr5tphr/zennit> | <https://github.com/virelay/corelay> | <https://github.com/virelay/virelay>
Accepted for publication at PLoS ONE
 21. Schwenk G and **Bach S** (2014).
“Detecting Behavioural and Structural Anomalies in Media-Cloud Applications”.
In: CoRR abs/1409.8035