# List of Publications

## Journal Articles

1. Weber L, Berend J, Weckbecker M, Binder A, Wiegand T, Samek W and **Lapuschkin S** (2025).
   "Efficient and Flexible Neural Network Training through Layer-wise Feedback Propagation".
   In: *Transactions on Machine Learning Research* 2835–8856.
   *https://github.com/leanderweber/layerwise-feedback-propagation*

2. Hedström A, Bommer P L, Burns T F, **Lapuschkin S**, Samek W and Höhne M-C M (2025).
   "Evaluating Interpretable Methods via Geometric Alignment of Functional Distortions".
   In: *Transactions on Machine Learning Research* 2835–8856.
   *https://github.com/annahedstroem/GEF/ | TMLR Survey Certification*

3. Bley F, **Lapuschkin S**, Samek W and Montavon G (2025).
   "Explaining Predictive Uncertainty by Exposing Second-Order Effects".
   In: *Pattern Recognition* 160:111171.
   *https://github.com/florianbley/XAI-2ndOrderUncertainty*

4. Vielhaben J, **Lapuschkin S**, Montavon G and Samek W (2024).
   "Explainable AI for Time Series via Virtual Inspection Layers".
   In: *Pattern Recognition* 150:110309.
   *https://github.com/jvielhaben/DFT-LRP*

5. Becker S, Vielhaben J, Ackermann M, Müller K-R, **Lapuschkin S** and Samek W (2024).
   "AudioMNIST: Exploring Explainable Artificial Intelligence for Audio Analysis on a Simple Benchmark".
   In: *Journal of the Franklin Institute* 361(1):418–428.
   *https://github.com/soerenab/AudioMNIST*

6. Achtibat R, Dreyer M, Eisenbraun I, Bosse S, Wiegand T, Samek W and **Lapuschkin S** (2023).
   "From attribution maps to human-understandable explanations through Concept Relevance Propagation".
   In: *Nature Machine Intelligence* 5(9):1006–1019.
   *https://github.com/rachtibat/zennit-crp | https://github.com/maxdreyer/crp-human-study*

7. Hedström A, Bommer P, Wickstrøm K K, Samek W, **Lapuschkin S** and Höhne M-C M (2023).
   "The Meta-Evaluation Problem in Explainable AI: Identifying Reliable Estimators with MetaQuantus".
   In: *Transactions on Machine Learning Research* 2835–8856.
   *https://github.com/annahedstroem/MetaQuantus*

8. Weber L, **Lapuschkin S**, Binder A and Samek W (2023).
   "Beyond Explaining: Opportunities and Challenges of XAI-Based Model Improvement".
   In: *Information Fusion* 92:154–176

9. Hedström A, Weber L, Krakowczyk D G, Bareeva D, Motzkus F, Samek W, **Lapuschkin S** and Höhne M-C M (2023).
   "Quantus: An Explainable AI Toolkit for Responsible Evaluation of Neural Network Explanations and Beyond".
   In: *Journal of Machine Learning Research* 24(34):1–11.
   *https://github.com/understandable-machine-intelligence-lab/quantus*

10. Hofmann S M, Beyer F, **Lapuschkin S**, Golterman O, Loeffler M, Müller K-R, Villringer A, Samek W and Witte A V (2022).
    "Towards the Interpretability of Deep Learning Models for Multi-modal Neuroimaging: Finding Structural Changes of the Ageing Brain".
    In: *NeuroImage* 261:119504

11. Ma J, Schneider L, **Lapuschkin S**, Achtibat R, Durchrau M, Krois J, Schwendicke F and Samek W (2022).
    "Towards Trustworthy AI in Dentistry".
    In: *Journal of Dental Research* 00220345221106086

12. Rieckmann A, Dworzynski P, Arras L, **Lapuschkin S**, Samek W, Onyebuchi A A, Rod N H, Ekstrøm C T (2022).
    "Causes of Outcome Learning: A Causal Inference-inspired Machine Learning Approach to Disentangling Common Combinations of Potential Causes of a Health Outcome".
    In: *International Journal of Epidemiology* dyac078.
    *https://github.com/ekstroem/cool | https://www.causesofoutcomelearning.org*

13. Slijepcevic D, Horst F, **Lapuschkin S**, Horsak B, Raberger A-M, Kranzl A, Samek W, Breiteneder C, Schöllhorn W I and Zeppelzauer M (2022).

"Explaining Machine Learning Models for Clinical Gait Analysis".
In: *ACM Transactions on Computing for Healthcare* 3(2):14:1–27.
*https://github.com/sebastian-lapuschkin/explaining-deep-clinical-gait-classification*

14. Anders C J, Weber L, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2022).
"Finding and Removing Clever Hans: Using Explanation Methods to Debug and Improve Deep Models".
In: *Information Fusion* 77:261–295

15. Sun J, **Lapuschkin S**, Samek W and Binder A (2022).
"Explain and Improve: LRP-inference Fine-tuning for Image Captioning Models".
In: *Information Fusion* 77:233–246

16. Samek W, Montavon G, **Lapuschkin S**, Anders C J, and Müller K-R (2021).
"Explaining Deep Neural Networks and Beyond: A Review of Methods and Applications".
In: *Proceedings of the IEEE* 109(3):247–278

17. Yeom S-K, Seegerer P, **Lapuschkin S**, Binder A, Wiedemann S, Müller K-R and Samek W (2021).
"Pruning by Explaining: A Novel Criterion for Deep Neural Network Pruning".
In: *Pattern Recognition* 115:107899.
*https://github.com/seulkiyeom/LRP_pruning | https://github.com/seulkiyeom/LRP_Pruning_toy_example*

18. Aeles J, Horst F, **Lapuschkin S**, Lacourpaille L, and Hug F (2021).
"Revealing the Unique Features of Each Individual's Muscle Activation Signatures".
In: *Journal of the Royal Society Interface* 18(174):20200770.
*https://github.com/sebastian-lapuschkin/interpretable-emg-signatures*

19. Horst F, Slijepcevic D, Zeppelzauer M, Raberger AM, **Lapuschkin S**, Samek W, Schöllhorn WI, Breiteneder C, and Horsak B (2020).
"Explaining Automated Gender Classification of Human Gait".
In: *Gait & Posture* 81(S1):159–160

20. Hägele M, Seegerer P, **Lapuschkin S**, Bockmayr M, Samek W, Klauschen F, Müller K-R and Binder A (2020).
"Resolving Challenges in Deep Learning-based Analyses of Histopathological Images using Explanation Methods".
In: *Scientific Reports* 10:6423

21. Alber M, **Lapuschkin S**, Seegerer P, Hägele M, Schütt K T, Montavon G, Samek W, Müller K-R, Dähne S and Kindermans P-J (2019).
"iNNvestigate Neural Networks!".
In: *Journal of Machine Learning Research* 20(93):1–8.
*https://github.com/albermax/innvestigate*

22. **Lapuschkin S**, Wäldchen S, Binder A, Montavon G, Samek W and Müller K-R (2019).
"Unmasking Clever Hans Predictors and Assessing what Machines Really Learn".
In: *Nature Communications* 10:1069

23. Horst F, **Lapuschkin S**, Samek W, Müller K-R and Schöllhorn W I (2019).
"Explaining the Unique Nature of Individual Gait Patterns with Deep Learning".
In: *Scientific Reports* 9:2391.
*https://github.com/sebastian-lapuschkin/interpretable-deep-gait*

24. Montavon G, **Lapuschkin S**, Binder A, Samek W and Müller K-R (2017).
"Explaining NonLinear Classification Decisions with Deep Taylor Decomposition".
In: *Pattern Recognition* 65:211–222.
*Pattern Recognition Best Paper Award and Pattern Recognition Medal winner*

25. Samek W, Binder A, Montavon G, **Lapuschkin S**, and Müller K-R (2017).
"Evaluating the Visualization of what a Deep Neural Network has Learned".
In: *IEEE Transactions of Neural Networks and Learning Systems*

26. Sturm I, **Lapuschkin S**, Samek W and Müller K-R (2016).
"Interpretable Deep Neural Networks for Single-Trial EEG Classification".
In: *Journal of Neuroscience Methods* 274:141–145

27. **Lapuschkin S**, Binder A, Montavon G, Müller K-R and Samek W (2016).
"The Layer-wise Relevance Propagation Toolbox for Artificial Neural Networks".
In: *Journal of Machine Learning Research* 17(114):1–5.
*https://github.com/sebastian-lapuschkin/lrp_toolbox*

28. **Bach S**, Binder A, Montavon G, Klauschen F, Müller K-R and Samek W (2015).
"On Pixel-wise Explanations for Non-Linear Classifier Decisions by Layer-wise Relevance Propagation".
In: *PLoS ONE* 10(7):e0130140

## Contributions to Conference Proceedings and Workshops

1. Pahde F, Dreyer M, Weckbecker M, Weber L, Anders C J, Wiegand T, Samek W and **Lapuschkin S** (2025).
   "Navigating Neural Space: Revisiting Concept Activation Vectors to Overcome Directional Divergence".
   In: *Proceedings of the International Conference on Learning Representations* (*ICLR*) TBA.
   *https://github.com/frederikpahde/pattern-cav*

2. Bareeva D, Yolcu GÜ, Hedström A, Wiegand T, Samek W **Lapuschkin S** (2024).
   "Quanda: An Interpretability Toolkit for Training Data Attribution Evaluation and Beyond".
   In: *NeuRIPS 2024 Workshop on Attributing Model Behavior at Scale* (*ATTRIB 2024*) .
   *https://github.com/dilyabareeva/quanda*

3. Naujoks J R, Krasowski A, Weckbecker M, Wiegand T, **Lapuschkin S**, Samek W and Klausen R P (2024).
   "PINNfluence: Influence Functions for Physics-Informed Neural Networks".
   In: *NeuRIPS 2024 Workshop on Machine Learning and the Physical Sciences* (*ML4PS*) .
   *https://github.com/aleks-krasowski/PINNfluence*
   *Reproducibility Badge Winner*

4. Kopf L, Bommer P L, Hedström A, **Lapuschkin S**, Höhne M M-C and Bykov K (2024).
   "CoSy: Evaluating Textual Explanations of Neurons".
   In: *Advances in Neural Information Processing Systems* (*NeuRIPS*) 34656–34685. (*OpenReview*)
   *https://github.com/lkopf/cosy*

5. Nobis G, Springenberg M, Aversa M, Detzel M, Daems R, Murray-Smith R, Nakajima S, **Lapuschkin S**, Ermon S, Birdal T, Opper M, Knochenhauer C, Oala L and Samek W (2024).
   "Generative Fractional Diffusion Models".
   In: *Advances in Neural Information Processing Systems* (*NeuRIPS*) 25469–25509. (*OpenReview*)
   *https://github.com/GabrielNobis/gfdm*

6. Mekala R R, Pahde F, Baur S, Chandrashekar S, Diep M, Wenzel M A, Wisotzky E L, Yolcu G Ü, **Lapuschkin S**, Ma J, Eisert P, Lindvall M, Porter A and Samek W (2024).
   "Synthetic Generation of Dermatoscopic Images with GAN and Closed-Form Factorization".
   In: *ECCV 2024 Workshop on Synthetic Data for Computer Vision* (*SyntheticData4CV*) TBA. (*Green Open Access*)

7. Achtibat R, Hatefi S M V, Dreyer M, Jain A, Wiegand T, **Lapuschkin S**, Samek W (2024).
   "AttnLRP: Attention-Aware Layer-wise Relevance Propagation for Transformers".
   In: *Proceedings of the 41st International Conference on Machine Learning* (*ICML*) 135–168.
   *https://github.com/rachtibat/LRP-for-Transformers*

8. Hatefi S M V, Dreyer M, Achtibat R, Wiegand T, Samek W and **Lapuschkin S** (2024).
   "Pruning By Explaining Revisited: Optimizing Attribution Methods to Prune CNNs and Transformers".
   In: *Proceedings of the European Conference on Computer Vision* (*ECCV*) *Workshops* 152–169. (*Green Open Access*)
   *https://github.com/erfanhatefi/Pruning-by-eXplaining-in-PyTorch*

9. Hedström A, Weber L, **Lapuschkin S**, Höhne M M-C (2024).
   "A Fresh Look at Sanity Checks for Saliency Maps".
   In: *Proceedings of the 2nd XAI World Conference* 403–420. (*Green Open Access*)
   *https://github.com/annahedstroem/sanity-checks-revisited*

10. Tinauer C, Damulina A, Sackl M, Soellradl M, Achtibat R, Dreyer M, Pahde F, **Lapuschkin S**, Schmidt R, Ropele S, Samek W, Langkammer C (2024).
    "Explainable Concept Mappings of MRI: Revealing the Mechanisms Underlying Deep Learning-based Brain Disease Classification".
    In: *Proceedings of the 2nd XAI World Conference* 202–216. (*Green Open Access*)

11. Dreyer M, Purelku E, Vielhaben J, Samek W, **Lapuschkin S** (2024).
    "PURE: Turning Polysemantic Neurons Into Pure Features by Identifying Relevant Circuits".
    In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (*CVPR*) *Workshops* 8212–8217.
    *https://github.com/maxdreyer/PURE | Spotlight Paper*

12. Bareeva D, Dreyer M, Pahde F, Samek W and **Lapuschkin S** (2024).
    "Reactive Model Correction: Mitigating Harm to Task-Relevant Features via Conditional Bias Suppression".
    In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (*CVPR*) *Workshops* 3532–3541.
    *https://github.com/dilyabareeva/reactive_correction*

13. Dreyer M, Achtibat R, Samek W and **Lapuschkin S** (2024).
    "Understanding the (Extra-)Ordinary: Validating Deep Model Decisions with Prototypical Concept-based Explanations".
    In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (*CVPR*) *Workshops* 3491–3501.
    *https://github.com/maxdreyer/pcx*

14. Dreyer M, Pahde F, Anders C J, Samek W and **Lapuschkin S** (2024).
    "From Hope to Safety: Unlearning Biases of Deep Models via Gradient Penalization in Latent Space".
    In: *Proceedings of the AAAI Conference on Artificial Intelligence* (*AAAI*) 38(19):21046–21054.
    *https://github.com/frederikpahde/rrclarc*

15. Dawoud K, Samek W, Eisert P, **Lapuschkin S** and Bosse S (2023).
    "Human-Centered Evaluation of XAI Methods".
    In: *Proceedings of the IEEE International Conference on Data Mining* (*ICDM*) 912–921. (*Green Open Access*)

16. Frommholz A, Seipel F, **Lapuschkin S**, Samek W and Vielhaben J (2023).
    "XAI-based Comparison of Audio Event Classifiers with different Input Representations".
    In: *Proceedings of the International Conference on Content-based Multimedia Indexing* (*CBMI*) 126–132

17. Hedström A, Weber L, **Lapuschkin S** and Höhne M M-C (2023).
    "Sanity Checks Revisited: An Exploration to Repair the Model Parameter Randomisation Test".
    In: *NeuRIPS 2023 Workshop on XAIX* (*XAI in Action: Past, Present, and Future Applications*) (vVpefYmnsG)

18. Pahde F, Dreyer M, Samek W and **Lapuschkin S** (2023).
    "Reveal to Revise: An Explainable AI Life Cycle for Iterative Bias Correction of Deep Models".
    In: *Proceedings of the International Conference on Medical Image Computing and Computer-Assisted Intervention* 596–606. (*Green Open Access*)
    *https://github.com/maxdreyer/reveal2revise*

19. Binder A, Weber L, **Lapuschkin S**, Montavon G, Müller K-R and Samek W (2023).
    "Shortcomings of Top-Down Randomization-Based Sanity Checks for Evaluations of Deep Neural Network Explanations".
    In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (*CVPR*) 16143–16152

20. Dreyer M, Achtibat R, Wiegand T, Samek W and **Lapuschkin S** (2023).
    "Revealing Hidden Context Bias in Segmentation and Object Detection through Concept-specific Explanations".
    In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (*CVPR*) *Workshops* 3828–3838

21. Pahde F, Yolcu GÜ, Binder A, Samek W and **Lapuschkin S** (2023).
    "Optimizing Explanations by Network Canonization and Hyperparameter Search".
    In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (*CVPR*) *Workshops* 3818–3827

22. Krakowczyk D G, Prasse P, Reich D R, **Lapuschkin S**, Scheffer T, Jäger L A (2023).
    "Bridging the Gap: Gaze Events as Interpretable Concepts to Explain Deep Neural Sequence Models".
    In: *Proceedings of the Symposium on Eye Tracking Research and Applications* (*ETRA*) 1–8.
    *Best Short Paper Award Winner*

23. Krakowczyk D G, Reich D R, Prasse P, **Lapuschkin S**, Jäger L A and Scheffer T (2022).
    "Selection of XAI Methods Matters: Evaluation of Feature Attribution Methods for Oculomotoric Biometric Identification".
    In: *NeuRIPS 2022 Workshop on Gaze Meets ML* (GOLdDAP2AtI)

24. Motzkus F, Weber L and **Lapuschkin S** (2022).
    "Measurably Stronger Explanation Reliability via Model Canonization".
    In: *Proceedings of the International Conference on Image Processing* (*ICIP*) 516–520

25. Ede S, Baghdadlian S, Weber L, Nguyen A, Zanca D, Samek W and **Lapuschkin S** (2022).
    "Explain to Not Forget: Defending Against Catastrophic Forgetting with XAI".
    In: *Proceedings of the International Cross-Domain Conference for Machine Learning and Knowledge Extraction* (*CD-MAKE*) 1–18. (*Gold Open Access link*)

26. Sun J, **Lapuschkin S**, Samek W, Zhao Y, Cheung N-M and Binder A (2021).
    "Explanation-Guided Training for Cross-Domain Few-Shot Classification".
    In: *Proceedings of the 25th International Conference on Pattern Recognition* (*ICPR*) 7609–7616

27. Goh G S W, **Lapuschkin S**, Weber L, Samek W and Binder A (2021).
    "Understanding Integrated Gradients with SmoothTaylor for Deep Neural Network Attribution".
    In: *Proceedings of the 25th International Conference on Pattern Recognition* (*ICPR*) 4949–4956

28. Kohlbrenner M, Bauer A, Nakajima S, Binder A, Samek W, and **Lapuschkin S** (2020).
"Towards Best Practice in Explaining Neural Network Decisions with LRP".
In: *Proceedings of the IEEE International Joint Conference on Neural Networks (IJCNN)* 1-7

29. Sun J, **Lapuschkin S**, Samek W and Binder A (2020).
"Understanding Image Captioning Models beyond Visualizing Attention".
In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*

30. Anders C J, Neumann D, Marinč T, Samek W, Müller K-R and **Lapuschkin S** (2020).
"XAI for Analyzing and Unlearning Spurious Correlations in ImageNet".
In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*

31. Sun J, **Lapuschkin S**, Samek W, Zhao Y, Cheung N-M and Binder A (2020).
"Explain and Improve: Cross-Domain-Few-Shot-Learning Using Explanations".
In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*

32. Alber M, **Lapuschkin S**, Seegerer P, Hägele M, Schütt K T, Montavon G, Samek W, Müller K-R, Dähne S and Kindermans P-J (2018).
"How to iNNvestigate Neural Networks' Predictors!".
In: *Machine Learning Open Source Software: Sustainable Communities. NIPS Workshop*

33. **Lapuschkin S**, Binder A, Müller K-R and Samek W (2017).
"Understanding and Comparing Deep Neural Networks for Age and Gender Classification".
In: *Proceedings of the ICCV'17 Workshop on Analysis and Modeling of Faces and Gestures (AMFG)* 2017:1629-1638

34. Srinivasan V, **Lapuschkin S**, Hellge C, Müller K-R and Samek W (2017).
"Interpretable Action Recognition in Compressed Domain".
In: *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* 2017:1692-1696

35. **Bach S**, Binder A, Müller K-R and Samek W (2016).
"Controlling Explanatory Heatmap Resolution and Semantics via Decomposition Depth".
In: *Proceedings of the IEEE International Conference of Image Processing (ICIP)* 2016:2271-2275

36. Binder A, Samek W, Montavon G, **Bach S**, and Müller K-R (2016).
"Analyzing and Validating Neural Network Predictions".
In: *Proceedings of the ICML'16 Workshop on Visualization for Deep Learning* .
*Best Paper Award Winner*

37. **Lapuschkin S**, Binder A, Montavon G, Müller K-R and Samek W (2016).
"Analyzing Classifiers: Fisher Vectors and Deep Neural Networks".
In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* 2016:2912-2920

38. Montavon G, **Bach S**, Binder A, Samek W and Müller K-R (2016).
"Deep Taylor Decomposition of Neural Networks".
In: *Proceedings of the ICML'16 Workshop on Visualization for Deep Learning* 2016:1-3

39. Samek W, Montavon G, Binder A, **Lapuschkin S** and Müller K-R (2016).
"Interpreting the Predictions of Complex ML Models by Layer-wise Relevance Propagation".
In: *Proceedings of the Interpretable ML for Complex Systems NIPS'16 Workshop*

**Books**

1. Longo L, **Lapuschkin S** and Seifert C, editors (2024).
"Explainable Artificial Intelligence (Second World Conference, xAI 2024, Valletta, Malta, July 17–19, 2024, Proceedings, Part I-IV)".
Springer (Cham), Part I ISBN: 978-3-031-63787-2. *Part II ISBN: 978-3-031-63797-1.*
*Part III ISBN: 978-3-031-63800-8. Part IV ISBN: 978-3-031-63803-9*

**Book Chapters**

1. Becking D, Dreyer M, Samek W, Müller K and **Lapuschkin S** (2022).
"ECQ$^x$: Explainability-Driven Quantization for Low-Bit and Sparse DNNs".
In: *xxAI – Beyond Explainable AI* 271-296. Springer, Cham

2. Montavon G, Binder A, **Lapuschkin S**, Samek W and Müller K-R (2019).
"Layer-wise relevance propagation: An Overview".
In: *Explainable AI: Interpreting, Explaining and Visualizing Deep Learning* 193-209. Springer, Cham

3. Binder A, **Bach S**, Montavon G, Müller K-R and Samek W (2016).
"Layer-wise Relevance Propagation for Deep Neural Network Architectures".
In: *Information Science and Applications (ICISA) 2016. Lecture Notes in Electrical Engineering* 276:913-922.
Springer, Singapore

4. Binder A, Montavon G, **Lapuschkin S**, Müller K-R and Samek W (2016).
   "Layer-wise Relevance Propagation for Neural Networks with Local Renormalization Layers".
   In: *Lecture Notes in Computer Science* 9887:63-71. Springer, Berlin/Heidelberg

**Preprints**

1. Labarta T, Hoang N, Weitz K, Samek W, **Lapuschkin S** and Weber L (2025).
   "See What I Mean? CUE: A Cognitive Model of Understanding Explanations".
   In: *CoRR abs/2506.14775*.
   *Accepted for publication at the IJCAI 2025 XAI Workshop*

2. Hatefi S M V, Dreyer M, Achtibat R, Kahardipraja P, Wiegand T, Samek W and **Lapuschkin S** (2025).
   "Attribution-guided Pruning for Compression, Circuit Discovery, and Targeted Correction in LLMs".
   In: *CoRR abs/2506.13727*.
   *https://github.com/erfanhatefi/SparC3*

3. Cantú E D, Wittmann R K, Abdeen O, Wagner P, Samek W, Baier M and **Lapuschkin S** (2025).
   "Deep Learning-based Multi Project InP Wafer Simulation for Unsupervised Surface Defect Detection".
   In: *CoRR abs/2506.10713*

4. Gururaj S, Grüne L, Samek W, **Lapuschkin S** and Weber L (2025).
   "Relevance-driven Input Dropout: an Explanation-guided Regularization Technique".
   In: *CoRR abs/2505.21595*.
   *https://github.com/Shreyas-Gururaj/LRP_Relevance_Dropout*

5. Dreyer M, Hufe L, Berend J, Wiegand T, **Lapuschkin S** and Samek W (2025).
   "From What to How: Attributing CLIP's Latent Components Reveals Unexpected Semantic Reliance".
   In: *CoRR abs/2505.20229*.
   *https://github.com/maxdreyer/attributing-clip*

6. Kahardipraja P, Achtibat R, Wiegand T, Samek W and **Lapuschkin S** (2025).
   "The Atlas of In-Context Learning: How Attention Heads Shape In-Context Retrieval Augmentation".
   In: *CoRR abs/2505.15807*. *https://github.com/pkhdipraja/in-context-atlas*

7. Hufe L, Venhoff C, Dreyer M, **Lapuschkin S** and Samek W (2025).
   "Dyslexify: A Mechanistic Defense Against Typographic Attacks in CLIP".
   In: *OpenReview: yFPsLBa53p*

8. Hatefi S M V, Dreyer M, Achtibat R, Kahardipraja P, Wiegand T, Samek W and **Lapuschkin S** (2025).
   "Attribution-guided Pruning for Compression, Circuit Discovery, and Targeted Correction in LLMs".
   In: *OpenReview: YKF9pMAXRv*

9. Bareeva D, Höhne M M C, Warnecke A, Pirch L, Müller K-R, Rieck K, **Lapuschkin S** and Bykov K (2025).
   "Manipulating Feature Visualizations with Gradient Slingshots".
   In: *OpenReview: TgczQwE1Iu*

10. Zverev E, Kortukov E, Panfilov A, Volkova A, Tabesh S, **Lapuschkin S**, Samek W and Lampert C H (2025).
    "ASIDE: Architectural Separation of Instructions and Data in Language Models".
    In: *CoRR abs/2503.10566*

11. Joseph S, Suresh P, Hufe L, Stevinson E, Graham R, Vadi Y, Bzdok D, **Lapuschkin S**, Sharkey L and Richards B A (2025).
    "Prisma: An Open Source Toolkit for Mechanistic Interpretability in Vision and Video".
    In: *OpenReview: 2WpymqqWQm*.
    *https://huggingface.co/Prisma-Multimodal | https://github.com/soniajoseph/ViT-Prisma*
    *Accepted for publication at the MIV Workshop during CVPR 2025*

12. Erogullari E, **Lapuschkin S**, Samek W and Pahde F (2025).
    "Post-Hoc Concept Disentanglement: From Correlated to Isolated Concept Representations".
    In: *CoRR abs/2503.05522*.
    *https://github.com/erenerogullari/cav-disentanglement*
    *Accepted for publication at the XAI World Conference 2025*

13. Puri B, Jain A, Golimblevskaia E, Kahardipraja P, Wiegand T, Samek W and **Lapuschkin S** (2025).
    "FADE: Why Bad Descriptions Happen to Good Features".
    In: *CoRR abs/2502.16994*.
    *Accepted for publication with ACL Findings*
    *https://github.com/brunibrun/FADE*

14. Arras L, Puri B, Kahardipraja P, **Lapuschkin S** and Samek W (2025).
    "A Close Look at Decomposition-based XAI-Methods for Transformer Language Models".
    In: *CoRR abs/2502.15886*

15. Pahde F, Wiegand T, **Lapuschkin S** and Samek W (2025).
    "Ensuring Medical AI Safety: Explainable AI-Driven Detection and Mitigation of Spurious Model Behavior and Associated Data".
    In: *CoRR abs/2501.13818.*
    *https://github.com/frederikpahde/medical-ai-safety*

16. Dreyer M, Berend J, Labarta T, Vielhaben J, Wiegand T, **Lapuschkin S** and Samek W (2025).
    "Mechanistic understanding and validation of large AI models with SemanticLens".
    In: *CoRR abs/2501.05398.*
    *Accepted for publication in Nature Machine Intelligence*
    *https://github.com/jim-berend/semanticlens | Demo: https://semanticlens.hhi-research-insights.eu/umap-view*

17. Yolcu G Ü, Wiegand T, Samek W and **Lapuschkin S** (2024).
    "DualView: Data Attribution from the Dual Perspective".
    In: *CoRR abs/2402.12118.*
    *https://github.com/gumityolcu/DualView*

18. Gerstenberger M, **Lapuschkin S**, Eisert P and Bosse S (2022).
    "But That's Not Why: Inference Adjustment by Interactive Prototype Deselection".
    In: *CoRR abs/2203.10087*

19. Anders C J, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2021).
    "Software for Dataset-wide XAI: From Local Explanations to Global Insights with Zennit, CoRelAy, and ViRelAy".
    In: *CoRR abs/2106.13200. https://github.com/chr5tphr/zennit |*
    *https://github.com/virelay/corelay | https://github.com/virelay/virelay*

20. Schwenk G and **Bach S** (2014).
    "Detecting Behavioural and Structural Anomalies in Media-Cloud Applications".
    In: *CoRR abs/1409.8035*