


Publications

Summary of Scientific Impact

	All	Since 2021
# Publications	109	83
# Citations	23645	19630
h-index	40	39
i10-index	68	66



per Google Scholar, retrieved on May 30th, 2026.

List of Publications

Journal Articles

1. Anders C J, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2026).
“Software for Dataset-wide XAI: From Local Explanations to Global Insights with Zennit, CoRelAy, and ViRelAy”.
In: *PLoS ONE* 21(1):e0336683.
<https://github.com/chr5tphr/zennit> | <https://github.com/virelay/corelay> | <https://github.com/virelay/virelay>
2. Cantú E D, Wittmann R K, Abdeen O, Wagner P, Samek W, Baier M and **Lapuschkin S** (2026).
“Deep Learning-based Multi Project InP Wafer Simulation towards Unsupervised Surface Defect Detection”.
In: *Journal of the Franklin Institute* 363(2):108316
3. Yolcu G Ü, Weckbecker M, Wiegand T, Samek W and **Lapuschkin S** (2025).
“Sparse, Efficient and Explainable Data Attribution with DualXDA”.
In: *Transactions on Machine Learning Research* qfx81N884A.
<https://github.com/gumityolcu/DualXDA>
4. Dreyer M, Berend J, Labarta T, Vielhaben J, Wiegand T, **Lapuschkin S** and Samek W (2025).
“Mechanistic Understanding and Validation of Large AI Models with SemanticLens”.
In: *Nature Machine Intelligence* 1–14.
<https://github.com/jim-berend/semanticlens> | Demo: <https://semanticlens.hhi-research-insights.eu>
5. Pahde F, Wiegand T, **Lapuschkin S** and Samek W (2025).
“Ensuring Medical AI Safety: Explainable AI-Driven Detection and Mitigation of Spurious Model Behavior and Associated Data”.
In: *Machine Learning* 114(9):206.
<https://github.com/frederikpahde/medical-ai-safety>
6. Ma J, Weicken E, Pahde F, Weitz K, **Lapuschkin S**, Samek W and Wiegand T (2025).
“Künstliche Intelligenz auf dem Prüfstand: Anforderungen, Qualitätskriterien und Prüfwerkzeuge für medizinische Anwendungen [Artificial intelligence under scrutiny: requirements, quality criteria, and testing tools for medical applications]”.
In: *Bundesgesundheitsblatt – Gesundheitsforschung – Gesundheitsschutz* 68:915-923
7. Storås A M, Dreyer M, Pahde F, **Lapuschkin S**, Samek W, Halvorsen P, de Lange T, Mori Y, Hann A, Berzin T M, Parasa S and Riegler M A (2025).
“Exploring the Clinical Value of Concept-based AI Explanations in Gastrointestinal Disease Detection”.
In: *Scientific Reports* 15(1):28860.
<https://github.com/AndreaStoraas/conceptXAI-GItract>
8. Weber L, Berend J, Weckbecker M, Binder A, Wiegand T, Samek W and **Lapuschkin S** (2025).
“Efficient and Flexible Neural Network Training through Layer-wise Feedback Propagation”.
In: *Transactions on Machine Learning Research* 9oToxYVOSW.
<https://github.com/leanderweber/layerwise-feedback-propagation>
9. Hedström A, Bommer P L, Burns T F, **Lapuschkin S**, Samek W and Höhne M-C M (2025).
“Evaluating Interpretable Methods via Geometric Alignment of Functional Distortions”.
In: *Transactions on Machine Learning Research* ukLxqA8zXj.
<https://github.com/annahedstroem/GEF/> | TMLR Survey Certification

10. Bley F, **Lapuschkin S**, Samek W and Montavon G (2025).
 “Explaining Predictive Uncertainty by Exposing Second-Order Effects”.
 In: *Pattern Recognition* 160:111171.
<https://github.com/florianbley/XAI-2ndOrderUncertainty>
11. Vielhaben J, **Lapuschkin S**, Montavon G and Samek W (2024).
 “Explainable AI for Time Series via Virtual Inspection Layers”.
 In: *Pattern Recognition* 150:110309.
<https://github.com/jvielhaben/DFT-LRP>
12. Becker S, Vielhaben J, Ackermann M, Müller K-R, **Lapuschkin S** and Samek W (2024).
 “AudioMNIST: Exploring Explainable Artificial Intelligence for Audio Analysis on a Simple Benchmark”.
 In: *Journal of the Franklin Institute* 361(1):418–428.
<https://github.com/soerenab/AudioMNIST>
13. Achibat R, Dreyer M, Eisenbraun I, Bosse S, Wiegand T, Samek W and **Lapuschkin S** (2023).
 “From attribution maps to human-understandable explanations through Concept Relevance Propagation”.
 In: *Nature Machine Intelligence* 5(9):1006–1019.
<https://github.com/rachtibat/zennit-crp> | <https://github.com/maxdreyer/crp-human-study>
14. Hedström A, Bommer P, Wickstrøm K K, Samek W, **Lapuschkin S** and Höhne M-C M (2023).
 “The Meta-Evaluation Problem in Explainable AI: Identifying Reliable Estimators with MetaQuantus”.
 In: *Transactions on Machine Learning Research* j3FK00HyfU.
<https://github.com/annahedstroem/MetaQuantus>
15. Weber L, **Lapuschkin S**, Binder A and Samek W (2023).
 “Beyond Explaining: Opportunities and Challenges of XAI-Based Model Improvement”.
 In: *Information Fusion* 92:154–176
16. Hedström A, Weber L, Krakowczyk D G, Bareeva D, Motzkus F, Samek W, **Lapuschkin S** and Höhne M-C M (2023).
 “Quantus: An Explainable AI Toolkit for Responsible Evaluation of Neural Network Explanations and Beyond”.
 In: *Journal of Machine Learning Research* 24(34):1–11.
<https://github.com/understandable-machine-intelligence-lab/quantus>
17. Hofmann S M, Beyer F, **Lapuschkin S**, Goltermann O, Loeffler M, Müller K-R, Villringer A, Samek W and Witte A V (2022).
 “Towards the Interpretability of Deep Learning Models for Multi-modal Neuroimaging: Finding Structural Changes of the Ageing Brain”.
 In: *NeuroImage* 261:119504
18. Ma J, Schneider L, **Lapuschkin S**, Achibat R, Duchrau M, Krois J, Schwendicke F and Samek W (2022).
 “Towards Trustworthy AI in Dentistry”.
 In: *Journal of Dental Research* 00220345221106086
19. Rieckmann A, Dworzynski P, Arras L, **Lapuschkin S**, Samek W, Onyebuchi A A, Rod N H, Ekstrøm C T (2022).
 “Causes of Outcome Learning: A Causal Inference-inspired Machine Learning Approach to Disentangling Common Combinations of Potential Causes of a Health Outcome”.
 In: *International Journal of Epidemiology* dyac078.
<https://github.com/ekstroem/cool> | <https://www.causesofoutcomelearning.org>
20. Slijepcevic D, Horst F, **Lapuschkin S**, Horsak B, Raberger A-M, Kranzl A, Samek W, Breiteneder C, Schöllhorn W I and Zeppelzauer M (2022).
 “Explaining Machine Learning Models for Clinical Gait Analysis”.
 In: *ACM Transactions on Computing for Healthcare* 3(2):14:1–27.
<https://github.com/sebastian-lapuschkin/explaining-deep-clinical-gait-classification>
21. Anders C J, Weber L, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2022).
 “Finding and Removing Clever Hans: Using Explanation Methods to Debug and Improve Deep Models”.
 In: *Information Fusion* 77:261–295
22. Sun J, **Lapuschkin S**, Samek W and Binder A (2022).
 “Explain and Improve: LRP-inference Fine-tuning for Image Captioning Models”.
 In: *Information Fusion* 77:233–246
23. Samek W, Montavon G, **Lapuschkin S**, Anders C J, and Müller K-R (2021).
 “Explaining Deep Neural Networks and Beyond: A Review of Methods and Applications”.
 In: *Proceedings of the IEEE* 109(3):247–278

24. Yeom S-K, Seegerer P, **Lapuschkin S**, Binder A, Wiedemann S, Müller K-R and Samek W (2021).
 “Pruning by Explaining: A Novel Criterion for Deep Neural Network Pruning”.
 In: *Pattern Recognition* 115:107899.
https://github.com/seulkiyeom/LRP_pruning | https://github.com/seulkiyeom/LRP_Pruning_toy_example
25. Aeles J, Horst F, **Lapuschkin S**, Lacourpaille L, and Hug F (2021).
 “Revealing the Unique Features of Each Individual’s Muscle Activation Signatures”.
 In: *Journal of the Royal Society Interface* 18(174):20200770.
<https://github.com/sebastian-lapuschkin/interpretable-emg-signatures>
26. Horst F, Slijepcevic D, Zeppelzauer M, Raberger AM, **Lapuschkin S**, Samek W, Schöllhorn WI, Breiteneder C, and Horsak B (2020).
 “Explaining Automated Gender Classification of Human Gait”.
 In: *Gait & Posture* 81(S1):159–160
27. Hägele M, Seegerer P, **Lapuschkin S**, Bockmayr M, Samek W, Klauschen F, Müller K-R and Binder A (2020).
 “Resolving Challenges in Deep Learning-based Analyses of Histopathological Images using Explanation Methods”.
 In: *Scientific Reports* 10:6423
28. Alber M, **Lapuschkin S**, Seegerer P, Hägele M, Schütt K T, Montavon G, Samek W, Müller K-R, Dähne S and Kindermans P-J (2019).
 “iNNvestigate Neural Networks!”.
 In: *Journal of Machine Learning Research* 20(93):1–8.
<https://github.com/albermax/innvestigate>
29. **Lapuschkin S**, Wäldchen S, Binder A, Montavon G, Samek W and Müller K-R (2019).
 “Unmasking Clever Hans Predictors and Assessing what Machines Really Learn”.
 In: *Nature Communications* 10:1069
30. Horst F, **Lapuschkin S**, Samek W, Müller K-R and Schöllhorn W I (2019).
 “Explaining the Unique Nature of Individual Gait Patterns with Deep Learning”.
 In: *Scientific Reports* 9:2391.
<https://github.com/sebastian-lapuschkin/interpretable-deep-gait>
31. Montavon G, **Lapuschkin S**, Binder A, Samek W and Müller K-R (2017).
 “Explaining NonLinear Classification Decisions with Deep Taylor Decomposition”.
 In: *Pattern Recognition* 65:211–222.
Pattern Recognition Best Paper Award and Pattern Recognition Medal winner
32. Samek W, Binder A, Montavon G, **Lapuschkin S**, and Müller K-R (2017).
 “Evaluating the Visualization of what a Deep Neural Network has Learned”.
 In: *IEEE Transactions of Neural Networks and Learning Systems*
33. Sturm I, **Lapuschkin S**, Samek W and Müller K-R (2016).
 “Interpretable Deep Neural Networks for Single-Trial EEG Classification”.
 In: *Journal of Neuroscience Methods* 274:141–145
34. **Lapuschkin S**, Binder A, Montavon G, Müller K-R and Samek W (2016).
 “The Layer-wise Relevance Propagation Toolbox for Artificial Neural Networks”.
 In: *Journal of Machine Learning Research* 17(114):1–5.
https://github.com/sebastian-lapuschkin/lrp_toolbox
35. **Bach S**, Binder A, Montavon G, Klauschen F, Müller K-R and Samek W (2015).
 “On Pixel-wise Explanations for Non-Linear Classifier Decisions by Layer-wise Relevance Propagation”.
 In: *PLoS ONE* 10(7):e0130140

Contributions to Conference Proceedings and Workshops

1. Ostermann S, Gurgurov D, Baeumel T, Hedderich M A, **Lapuschkin S**, Samek W and Schmitt V (2026).
 “From Weights to Activations: Is Steering the Next Frontier of Adaptation?”.
 In: *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*
2. Komorowski P, Golimblevskaia E, Achtibat R, Wiegand T, **Lapuschkin S** and Samek W (2026).
 “Attribution-Guided Decoding”.
 In: *Proceedings of the International Conference on Learning Representations (ICLR)* .
<https://github.com/piotr-komorowski/attribution-guided-decoding>
3. Kahardipraja P, Achtibat R, Wiegand T, Samek W and **Lapuschkin S** (2025).
 “The Atlas of In-Context Learning: How Attention Heads Shape In-Context Retrieval Augmentation”.
 In: *Advances in Neural Information Processing Systems (NeurIPS)* .
<https://github.com/pkhardipraja/in-context-atlas>

4. Bareeva D, Höhne M M C, Warnecke A, Pirch L, Müller K-R, Rieck K, **Lapuschkin S** and Bykov K (2025).
“Manipulating Feature Visualizations with Gradient Slingshots”.
In: *Advances in Neural Information Processing Systems (NeurIPS)* .
https://github.com/dilyabareeva/grad_slingshot
5. Sandmann E, **Lapuschkin S** and Samek W (2025).
“Iterative Inference in a Chess-Playing Neural Network”.
In: *NeurIPS 2025 Workshop on Mechanistic Interpretability* .
<https://github.com/hartigel/leela-logit-lens>
6. Golimblevskaia E, Jain A, Puri B, Ibrahim A, Samek W and **Lapuschkin S** (2025).
“Circuit Insights: Towards Interpretability Beyond Activations”.
In: *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP (EMNLP)* .
<https://github.com/egolimblevskaia/WeightLens> | <https://github.com/egolimblevskaia/CircuitLens>
7. Labarta T, Hoang N, Weitz K, Samek W, **Lapuschkin S** and Weber L (2025).
“See What I Mean? CUE: A Cognitive Model of Understanding Explanations”.
In: *Proceedings of the IJCAI Workshops 2025: XAI Workshop* .
<https://arxiv.org/abs/2506.14775>
8. Puri B, Jain A, Golimblevskaia E, Kahardipraja P, Wiegand T, Samek W and **Lapuschkin S** (2025).
“FADE: Why Bad Descriptions Happen to Good Features”.
In: *Findings of the Association for Computational Linguistics (ACL)* 17138–17160.
<https://github.com/brunibrun/FADE>
9. Naujoks J, Krasowski A, Weckbecker M, Yolcu G Ü, Wiegand T, **Lapuschkin S**, Samek W and Klausen R P (2025).
“Leveraging Influence Functions for Resampling Data in Physics-Informed Neural Networks”.
In: *Proceedings of the 3rd XAI World Conference* 383–395.
https://github.com/aleks-krasowski/PINNfluence_resampling
10. Erogullari E, **Lapuschkin S**, Samek W and Pahde F (2025).
“Post-Hoc Concept Disentanglement: From Correlated to Isolated Concept Representations”.
In: *Proceedings of the 3rd XAI World Conference* 68–89.
<https://github.com/erenerogullari/cav-disentanglement>
11. Joseph S, Suresh P, Hufe L, Stevinson E, Graham R, Vadi Y, Bzdok D, **Lapuschkin S**, Sharkey L and Richards A (2025).
“Prisma: An Open Source Toolkit for Mechanistic Interpretability in Vision and Video”.
In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops: MIV Workshop* TBA.
<https://arxiv.org/abs/2504.19475> | <https://github.com/Prisma-Multimodal/ViT-Prisma>
12. Pahde F, Dreyer M, Weckbecker M, Weber L, Anders C J, Wiegand T, Samek W and **Lapuschkin S** (2025).
“Navigating Neural Space: Revisiting Concept Activation Vectors to Overcome Directional Divergence”.
In: *Proceedings of the International Conference on Learning Representations (ICLR)* .
<https://github.com/frederikpahde/pattern-cav>
13. Bareeva D, Yolcu GÜ, Hedström A, Wiegand T, Samek W and **Lapuschkin S** (2024).
“Quanda: An Interpretability Toolkit for Training Data Attribution Evaluation and Beyond”.
In: *NeurIPS 2024 Workshop on Attributing Model Behavior at Scale (ATTRIB 2024)* .
<https://github.com/dilyabareeva/quanda>
14. Naujoks J R, Krasowski A, Weckbecker M, Wiegand T, **Lapuschkin S**, Samek W and Klausen R P (2024).
“PINNfluence: Influence Functions for Physics-Informed Neural Networks”.
In: *NeurIPS 2024 Workshop on Machine Learning and the Physical Sciences (ML4PS)* .
<https://github.com/aleks-krasowski/PINNfluence>
Reproducibility Badge Winner
15. Kopf L, Bommer P L, Hedström A, **Lapuschkin S**, Höhne M M-C and Bykov K (2024).
“CoSy: Evaluating Textual Explanations of Neurons”.
In: *Advances in Neural Information Processing Systems (NeurIPS)* 34656–34685. (OpenReview)
<https://github.com/lkopf/cosy>
16. Nobis G, Springenberg M, Aversa M, Detzel M, Daems R, Murray-Smith R, Nakajima S, **Lapuschkin S**, Ermon S, Birdal T, Oppen M, Knochenhauer C, Oala L and Samek W (2024).
“Generative Fractional Diffusion Models”.
In: *Advances in Neural Information Processing Systems (NeurIPS)* 25469–25509. (OpenReview)
<https://github.com/GabrielNobis/gfdm>

17. Mekala R R, Pahde F, Baur S, Chandrashekar S, Diep M, Wenzel M A, Wisotzky E L, Yolcu G Ü, **Lapuschkin S**, Ma J, Eisert P, Lindvall M, Porter A and Samek W (2024).
 “Synthetic Generation of Dermatoscopic Images with GAN and Closed-Form Factorization”.
 In: *ECCV 2024 Workshop on Synthetic Data for Computer Vision (SyntheticData4CV)* 15642:368–384. (Green Open Access)
18. Achibat R, Hatefi S M V, Dreyer M, Jain A, Wiegand T, **Lapuschkin S**, Samek W (2024).
 “AttnLRP: Attention-Aware Layer-wise Relevance Propagation for Transformers”.
 In: *Proceedings of the 41st International Conference on Machine Learning (ICML)* 135–168.
<https://github.com/rachibat/LRP-for-Transformers>
19. Hatefi S M V, Dreyer M, Achibat R, Wiegand T, Samek W and **Lapuschkin S** (2024).
 “Pruning By Explaining Revisited: Optimizing Attribution Methods to Prune CNNs and Transformers”.
 In: *Proceedings of the European Conference on Computer Vision (ECCV) Workshops* 152–169. (Green Open Access)
<https://github.com/erfanhatefi/Pruning-by-eXplaining-in-PyTorch>
20. Hedström A, Weber L, **Lapuschkin S**, Höhne M M-C (2024).
 “A Fresh Look at Sanity Checks for Saliency Maps”.
 In: *Proceedings of the 2nd XAI World Conference* 403–420. (Green Open Access)
<https://github.com/annahedstroem/sanity-checks-revisited>
21. Tinauer C, Damulina A, Sackl M, Soellradl M, Achibat R, Dreyer M, Pahde F, **Lapuschkin S**, Schmidt R, Ropele S, Samek W, Langkammer C (2024).
 “Explainable Concept Mappings of MRI: Revealing the Mechanisms Underlying Deep Learning-based Brain Disease Classification”.
 In: *Proceedings of the 2nd XAI World Conference* 202–216. (Green Open Access)
22. Dreyer M, Pürelku E, Vielhaben J, Samek W, **Lapuschkin S** (2024).
 “PURE: Turning Polysemantic Neurons Into Pure Features by Identifying Relevant Circuits”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 8212–8217.
<https://github.com/maxdreyer/PURE|SpotlightPaper>
23. Bareeva D, Dreyer M, Pahde F, Samek W and **Lapuschkin S** (2024).
 “Reactive Model Correction: Mitigating Harm to Task-Relevant Features via Conditional Bias Suppression”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3532–3541.
https://github.com/dilyabareeva/reactive_correction
24. Dreyer M, Achibat R, Samek W and **Lapuschkin S** (2024).
 “Understanding the (Extra-)Ordinary: Validating Deep Model Decisions with Prototypical Concept-based Explanations”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3491–3501.
<https://github.com/maxdreyer/pcx>
25. Dreyer M, Pahde F, Anders C J, Samek W and **Lapuschkin S** (2024).
 “From Hope to Safety: Unlearning Biases of Deep Models via Gradient Penalization in Latent Space”.
 In: *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)* 38(19):21046–21054.
<https://github.com/frederikpahde/rrclarc>
26. Dawoud K, Samek W, Eisert P, **Lapuschkin S** and Bosse S (2023).
 “Human-Centered Evaluation of XAI Methods”.
 In: *Proceedings of the IEEE International Conference on Data Mining (ICDM)* 912–921. (Green Open Access)
27. Frommholz A, Seipel F, **Lapuschkin S**, Samek W and Vielhaben J (2023).
 “XAI-based Comparison of Audio Event Classifiers with different Input Representations”.
 In: *Proceedings of the International Conference on Content-based Multimedia Indexing (CBMI)* 126–132
28. Hedström A, Weber L, **Lapuschkin S** and Höhne M M-C (2023).
 “Sanity Checks Revisited: An Exploration to Repair the Model Parameter Randomisation Test”.
 In: *NeurIPS 2023 Workshop on XAI (XAI in Action: Past, Present, and Future Applications)* (vVpefYmnsG)
29. Pahde F, Dreyer M, Samek W and **Lapuschkin S** (2023).
 “Reveal to Revise: An Explainable AI Life Cycle for Iterative Bias Correction of Deep Models”.
 In: *Proceedings of the International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)* 596–606. (Green Open Access)
<https://github.com/maxdreyer/reveal2revise>

30. Binder A, Weber L, **Lapuschkin S**, Montavon G, Müller K-R and Samek W (2023).
 “Shortcomings of Top-Down Randomization-Based Sanity Checks for Evaluations of Deep Neural Network Explanations”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)* 16143–16152
31. Dreyer M, Achtabat R, Wiegand T, Samek W and **Lapuschkin S** (2023).
 “Revealing Hidden Context Bias in Segmentation and Object Detection through Concept-specific Explanations”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3828–3838
32. Pahde F, Yolcu GÜ, Binder A, Samek W and **Lapuschkin S** (2023).
 “Optimizing Explanations by Network Canonization and Hyperparameter Search”.
 In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops* 3818–3827
33. Krakowczyk D G, Prasse P, Reich D R, **Lapuschkin S**, Scheffer T, Jäger L A (2023).
 “Bridging the Gap: Gaze Events as Interpretable Concepts to Explain Deep Neural Sequence Models”.
 In: *Proceedings of the Symposium on Eye Tracking Research and Applications (ETRA)* 1–8.
Best Short Paper Award Winner
34. Krakowczyk D G, Reich D R, Prasse P, **Lapuschkin S**, Jäger L A and Scheffer T (2022).
 “Selection of XAI Methods Matters: Evaluation of Feature Attribution Methods for Oculomotoric Biometric Identification”.
 In: *NeurIPS 2022 Workshop on Gaze Meets ML (GOLDAP2AtI)*
35. Motzkus F, Weber L and **Lapuschkin S** (2022).
 “Measurably Stronger Explanation Reliability via Model Canonization”.
 In: *Proceedings of the International Conference on Image Processing (ICIP)* 516–520
36. Ede S, Baghdadlian S, Weber L, Nguyen A, Zanca D, Samek W and **Lapuschkin S** (2022).
 “Explain to Not Forget: Defending Against Catastrophic Forgetting with XAI”.
 In: *Proceedings of the International Cross-Domain Conference for Machine Learning and Knowledge Extraction (CD-MAKE)* 1–18. ([Gold Open Access link](#))
37. Sun J, **Lapuschkin S**, Samek W, Zhao Y, Cheung N-M and Binder A (2021).
 “Explanation-Guided Training for Cross-Domain Few-Shot Classification”.
 In: *Proceedings of the 25th International Conference on Pattern Recognition (ICPR)* 7609–7616
38. Goh G S W, **Lapuschkin S**, Weber L, Samek W and Binder A (2021).
 “Understanding Integrated Gradients with SmoothTaylor for Deep Neural Network Attribution”.
 In: *Proceedings of the 25th International Conference on Pattern Recognition (ICPR)* 4949–4956
39. Kohlbrenner M, Bauer A, Nakajima S, Binder A, Samek W, and **Lapuschkin S** (2020).
 “Towards Best Practice in Explaining Neural Network Decisions with LRP”.
 In: *Proceedings of the IEEE International Joint Conference on Neural Networks (IJCNN)* 1-7
40. Sun J, **Lapuschkin S**, Samek W and Binder A (2020).
 “Understanding Image Captioning Models beyond Visualizing Attention”.
 In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*
41. Anders C J, Neumann D, Marinč T, Samek W, Müller K-R and **Lapuschkin S** (2020).
 “XAI for Analyzing and Unlearning Spurious Correlations in ImageNet”.
 In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*
42. Sun J, **Lapuschkin S**, Samek W, Zhao Y, Cheung N-M and Binder A (2020).
 “Explain and Improve: Cross-Domain-Few-Shot-Learning Using Explanations”.
 In: *XXAI: Extending Explainable AI Beyond Deep Models and Classifiers. ICML Workshop*
43. Alber M, **Lapuschkin S**, Seegerer P, Hägele M, Schütt K T, Montavon G, Samek W, Müller K-R, Dähne S and Kindermans P-J (2018).
 “How to iNNvestigate Neural Networks’ Predictors!”.
 In: *Machine Learning Open Source Software: Sustainable Communities. NIPS Workshop*
44. **Lapuschkin S**, Binder A, Müller K-R and Samek W (2017).
 “Understanding and Comparing Deep Neural Networks for Age and Gender Classification”.
 In: *Proceedings of the ICCV’17 Workshop on Analysis and Modeling of Faces and Gestures (AMFG)* 2017:1629-1638
45. Srinivasan V, **Lapuschkin S**, Hellge C, Müller K-R and Samek W (2017).
 “Interpretable Action Recognition in Compressed Domain”.
 In: *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* 2017:1692-1696

46. **Bach S**, Binder A, Müller K-R and Samek W (2016).
 “Controlling Explanatory Heatmap Resolution and Semantics via Decomposition Depth”.
 In: *Proceedings of the IEEE International Conference of Image Processing (ICIP)* 2016:2271-2275
47. Binder A, Samek W, Montavon G, **Bach S**, and Müller K-R (2016).
 “Analyzing and Validating Neural Network Predictions”.
 In: *Proceedings of the ICML’16 Workshop on Visualization for Deep Learning* .
Best Paper Award Winner
48. **Lapuschkin S**, Binder A, Montavon G, Müller K-R and Samek W (2016).
 “Analyzing Classifiers: Fisher Vectors and Deep Neural Networks”.
 In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)* 2016:2912-2920
49. Montavon G, **Bach S**, Binder A, Samek W and Müller K-R (2016).
 “Deep Taylor Decomposition of Neural Networks”.
 In: *Proceedings of the ICML’16 Workshop on Visualization for Deep Learning* 2016:1-3
50. Samek W, Montavon G, Binder A, **Lapuschkin S** and Müller K-R (2016).
 “Interpreting the Predictions of Complex ML Models by Layer-wise Relevance Propagation”.
 In: *Proceedings of the Interpretable ML for Complex Systems NIPS’16 Workshop*

Books

1. Longo L, **Lapuschkin S** and Seifert C, editors (2024).
 “Explainable Artificial Intelligence (Second World Conference, xAI 2024, Valletta, Malta, July 17–19, 2024, Proceedings, Part I-IV)”.
 Springer (Cham), Part I ISBN: 978-3-031-63787-2. Part II ISBN: 978-3-031-63797-1.
 Part III ISBN: 978-3-031-63800-8. Part IV ISBN: 978-3-031-63803-9

Book Chapters

1. Becking D, Dreyer M, Samek W, Müller K and **Lapuschkin S** (2022).
 “ECQ^x: Explainability-Driven Quantization for Low-Bit and Sparse DNNs”.
 In: *xxAI – Beyond Explainable AI* 271-296. Springer, Cham
2. Montavon G, Binder A, **Lapuschkin S**, Samek W and Müller K-R (2019).
 “Layer-wise relevance propagation: An Overview”.
 In: *Explainable AI: Interpreting, Explaining and Visualizing Deep Learning* 193-209. Springer, Cham
3. Binder A, **Bach S**, Montavon G, Müller K-R and Samek W (2016).
 “Layer-wise Relevance Propagation for Deep Neural Network Architectures”.
 In: *Information Science and Applications (ICISA) 2016. Lecture Notes in Electrical Engineering* 276:913-922.
 Springer, Singapore
4. Binder A, Montavon G, **Lapuschkin S**, Müller K-R and Samek W (2016).
 “Layer-wise Relevance Propagation for Neural Networks with Local Renormalization Layers”.
 In: *Lecture Notes in Computer Science* 9887:63-71. Springer, Berlin/Heidelberg

Preprints

1. Zimmermann J P, Berend J, Loho G, **Lapuschkin S** and Samek W (2026).
 “Playing the Network Backward: A Game Theoretic Attribution Framework”.
 In: *CoRR abs/2605.06212*
2. Bouanani O, Berend J, Samek W, **Lapuschkin S** and Dreyer M (2026).
 “Contrastive Semantic Projection: Faithful Neuron Labeling with Contrastive Examples”.
 In: *CoRR abs/2604.22477*.
Accepted for publication at the 4th XAI World Conference (2026)
3. Labarta T, Dreyer M, Weitz K, Samek W and **Lapuschkin S** (2026).
 “From Attribution to Action: A Human-Centered Application of Activation Steering”.
 In: *CoRR abs/2604.11467*
4. Mesgari M M, Ma J, Samek W, **Lapuschkin S** and Weber L (2026).
 “Structural Compactness as a Complementary Criterion for Explanation Quality”.
 In: *CoRR abs/2603.29491*.
Accepted for publication at the 4th XAI World Conference (2026)
5. Heydari S, Said J, Yolcu G Ü, Kortukov E, Golimblevskaia E, Vlachos E, Mygdalis V, Pitas I, **Lapuschkin S** and Arras L (2026).
 “Concept-based Explanations of Segmentation and Detection Models in Natural Disaster Management”.
 In: *CoRR abs/2603.23020*.
Accepted as a late-breaking contribution at the 4th XAI World Conference (2026)

6. Krasowski A, Klausen R P, Celik A, **Lapuschkin S**, Samek W and Naujoks J (2026).
 “Building Trust in PINNs: Error Estimation through Finite Difference Methods”.
 In: *CoRR abs/2603.15526*.
Accepted as a late-breaking contribution at the 4th XAI World Conference (2026)
7. Labarta T, Hoang N, Dreyer M, Berend J, Hein O, Ma J, Samek W and **Lapuschkin S** (2026).
 “X-SYS: A Reference Architecture for Interactive Explanation Systems”.
 In: *CoRR abs/2602.12748*.
Accepted for publication at the 4th XAI World Conference (2026)
8. Puri B, Berend J, **Lapuschkin S** and Samek W (2025).
 “Atlas-Alignment: Making Interpretability Transferable Across Language Models”.
 In: *CoRR abs/2510.27413*
9. Klausen R P, Timofeev I, Frank J, Naujoks J, Wiegand T, **Lapuschkin S** and Samek W (2025).
 “LieSolver: A PDE-constrained solver for IBVPs using Lie symmetries”.
 In: *CoRR abs/2510.25731*
10. Panfilov A, Kortukov E, Nikolić K, Bethge M, **Lapuschkin S**, Samek W, Prabhu A, Andriushchenko M, Geiping J (2025).
 “Strategic Dishonesty Can Undermine AI Safety Evaluations of Frontier LLMs”.
 In: *CoRR abs/2509.18058*
11. Hufe L, Venhoff C, Pürelku E, Dreyer M, **Lapuschkin S** and Samek W (2025).
 “Dyslexify: A Mechanistic Defense Against Typographic Attacks in CLIP”.
 In: *CoRR abs/2508.20570*
12. Hatefi S M V, Dreyer M, Achibat R, Kahardipraja P, Wiegand T, Samek W, Binder A and **Lapuschkin S** (2025).
 “Attribution-Guided Pruning for Insight and Control: Circuit Discovery and Targeted Correction in Small-scale LLMs”.
 In: *CoRR abs/2506.13727*.
<https://github.com/erfanhatefi/SparC3>
13. Gururaj S, Grüne L, Samek W, **Lapuschkin S** and Weber L (2025).
 “Relevance-driven Input Dropout: an Explanation-guided Regularization Technique”.
 In: *CoRR abs/2505.21595*.
https://github.com/Shreyas-Gururaj/LRP_Relevance_Dropout
14. Dreyer M, Hufe L, Berend J, Wiegand T, **Lapuschkin S** and Samek W (2025).
 “From What to How: Attributing CLIP’s Latent Components Reveals Unexpected Semantic Reliance”.
 In: *CoRR abs/2505.20229*.
<https://github.com/maxdreyer/attributing-clip>
15. Zverev E, Kortukov E, Panfilov A, Volkova A, Tabesh S, **Lapuschkin S**, Samek W and Lampert C H (2025).
 “ASIDE: Architectural Separation of Instructions and Data in Language Models”.
 In: *CoRR abs/2503.10566*
16. Arras L, Puri B, Kahardipraja P, **Lapuschkin S** and Samek W (2025).
 “A Close Look at Decomposition-based XAI-Methods for Transformer Language Models”.
 In: *CoRR abs/2502.15886*
17. Gerstenberger M, **Lapuschkin S**, Eisert P and Bosse S (2022).
 “But That’s Not Why: Inference Adjustment by Interactive Prototype Deselection”.
 In: *CoRR abs/2203.10087*
18. Anders C J, Neumann D, Samek W, Müller K-R and **Lapuschkin S** (2021).
 “Software for Dataset-wide XAI: From Local Explanations to Global Insights with Zennit, CoRelAy, and ViRelAy”.
 In: *CoRR abs/2106.13200*. <https://github.com/chr5tphr/zennit> |
<https://github.com/virelay/corelay> | <https://github.com/virelay/virelay>
Accepted for publication at PLoS ONE
19. Schwenk G and **Bach S** (2014).
 “Detecting Behavioural and Structural Anomalies in Media-Cloud Applications”.
 In: *CoRR abs/1409.8035*